

The SageLang Programming Language

A Comprehensive Guide

Jacob Yates

July 2026

Contents

The SageLang Programming Language: A Comprehensive Guide	3
Executive Summary	3
Part 1: Language Overview and Design Philosophy	4
1.1 Design Goals and Target Use Case	4
1.2 Language Characteristics	4
1.3 Execution Model	5
1.4 Performance Characteristics	6
Part 2: Internal Architecture and Core Modules	6
2.1 Module Dependency Graph	6
2.2 Lexer (lexer.c / lexer.h)	6
2.3 Parser (parser.c)	7
2.4 Value System (value.h / value.c)	8
2.5 Environment and Scoping (env.h / env.c)	9
2.6 Garbage Collector (gc.h / gc.c)	10
2.7 Interpreter (interpreter.c / interpreter.h)	11
2.8 Module System (module.h / module.c)	13
Part 3: Language Features by Phase	15
Phase 1–3: Core Language (Arithmetic, Variables, Control Flow)	15
Phase 4: Garbage Collection	15
Phase 5–6: Data Structures and OOP	15
Phase 7: Advanced Control Flow, Exceptions, Generators	16
Phase 8: Modules	17
Phase 9: Low-Level Programming	17
Part 4: Writing SageLang Programs	20
4.1 Basic Syntax Patterns	20
4.2 Control Flow Patterns	21
4.3 Functions and Closures	22
4.4 Structs, Enums, Traits, and Classes	23
4.5 Exception Handling	25
4.6 Generators	25
4.7 Modules	26
4.8 Defer Statements	26
4.9 Pattern Matching (match/case/default)	27
Part 5: Implementation Details and Runtime Behavior	28
5.1 Memory Layout and Object Representation	28
5.2 Call Stack and Environments	28
5.3 Exception Propagation	28
5.4 Generator State Management	28

5.5 Operator Precedence and Associativity	29
5.6 Type Coercion and Truthiness	29
Part 6: Compilation and Execution	29
6.1 Building the Interpreter	29
6.2 Main Entry Point (main.c)	31
6.3 Program Execution Flow	33
6.4 Runtime Backends	34
Part 7: Advanced Topics and Internals	35
7.1 Memory Leaks and GC Limitations	35
7.2 Performance Characteristics	35
3 Current Recipe Benchmark	35
3 Sage vs Python 3 Benchmarks	36
7.3 Extending SageLang with Native Functions	37
7.4 Debugging and Introspection	37
7.5 Interpreter Safety Limits	37
Part 8: Future Directions and Design Notes	38
8.1 Planned Features (Incomplete Implementations)	38
8.2 Design Decisions and Rationale	38
Part 9: Bundled Library Modules (lib/)	39
9.1 Arrays Module (<code>lib/arrays.sage</code>)	39
9.2 Strings Module (<code>lib/strings.sage</code>)	39
9.3 Dicts Module (<code>lib/dicts.sage</code>)	40
9.4 Iter Module (<code>lib/iter.sage</code>)	40
9.5 Stats Module (<code>lib/stats.sage</code>)	41
9.6 Utils Module (<code>lib/utils.sage</code>)	41
9.7 Assert Module (<code>lib/assert.sage</code>)	42
9.8 Math Module (<code>lib/math.sage</code>)	42
9.9 GPU Libraries	43
9.10 OS Development Libraries	44
9.11 Networking Libraries	47
9.12 Cryptography Libraries	47
9.13 Machine Learning Libraries	47
9.14 CUDA Libraries	48
9.15 Standard Library (<code>lib/std/</code>)	49
9.16 LLM / Neural Network Libraries	50
9.17 Agent Framework (<code>lib/agent/</code>)	51
9.18 Chatbot Framework (<code>lib/chat/</code>)	52
9.19 UI Widget Library (<code>lib/graphics/ui.sage</code>)	52
9.20 AVR Assembler & Arduino Uno Support (<code>core/boards/AVR/</code>)	53
9.21 ESP32 Board Support (<code>core/boards/ESP32/</code> , <code>core/lib/esp32.sage</code>)	54
Part 13: Self-Hosting (Phase 13)	55
13.1 Architecture	55
13.2 Running Self-Hosted Code	55
13.3 Key Design Decisions	56
13.4 Feature Coverage	56
13.5 Test Suite	56
Part 10: Native Standard Library Modules (Phase 11)	57
10.1 Math Module	57
10.2 IO Module	57
10.3 String Module	58
10.4 Sys Module	58
10.5 FAT Module	58
Part 11: Concurrency and Async/Await (Phase 11)	59
11.1 Thread Module	59

11.2 Async/Await	59
11.3 GC Thread Safety	60
Part 12: Developer Tooling (Phase 12)	60
12.1 REPL (Read-Eval-Print Loop)	60
12.2 Code Formatter	61
12.3 Linter	61
12.4 OIS Package Management	62
12.5 Syntax Highlighting	62
12.6 Language Server Protocol (LSP)	62
12.7 CommonJS to ESM Transpiler (<code>cjs2esm</code>)	63
Example Programs	63
Part 14: Networking (Phase 14)	64
14.1 Socket Module	64
14.2 TCP Module	64
14.3 HTTP Module	65
14.4 SSL Module	65
14.5 Important Notes	66
Part 15: JSON Library (cJSON Port)	66
15.1 Basic Usage	66
15.2 Creating JSON	67
15.3 Sage-Native Conversion	67
15.4 API Reference	67
15.5 Important Notes	68
Conclusion	68
GPU Graphics (Vulkan + OpenGL)	68
Build Requirements	69
Quick Start: Empty Window	69
Quick Start: Triangle	70
GPU Module API Reference	70
Key Constants	73
Demos	73
9.21 Hardware Natives (Embedded Targets)	74
Appendix: Quick Reference	75
Keywords	75
Built-in Functions	75
Operators	76

The SageLang Programming Language: A Comprehensive Guide

Executive Summary

SageLang is a **Python-inspired, systems-oriented programming language** written in C. It combines familiar Python syntax (indentation-based blocks, dynamic typing) with low-level systems capabilities (garbage collection, exception handling, generators, and module imports). The language now supports ten execution backends (C, LLVM IR, native assembly, bytecode VM, SageMetal VM, JIT, AOT, Kotlin/Android) and a self-hosted interpreter written in Sage itself. As of v3, Sage features structural value equality in uniqueness checks, safe non-hanging string/value repeating, and robust tab/whitespace token checks in sandbox security guards. This guide documents the language design, internal architecture, runtime semantics, and practical usage patterns derived from the complete C source implementation.

Part 1: Language Overview and Design Philosophy

Note on PDF Generation: To generate the official PDF manual (`The_Sage_Programming_Language.pdf`), install `pandoc` and `texlive-xetex`, then execute `cd core/docs && pandoc SageLang_Guide.md -o The_Sage_Programming_Language.pdf --pdf-engine=xelatex -V geometry:margin=1in`.

1.1 Design Goals and Target Use Case

SageLang supports execution on various platforms including `x86_64`, `aarch64`, Windows, macOS, Linux, and microcontrollers such as the classic ESP32 (verified on ESP32-D0WD-V3 / 4MB flash). The `core/lib/esp32.sage` module and `core/boards/ESP32/` package provide full ESP32 support.

SageLang is designed as an **educational and practical embedded scripting language** that:

- **Bridges Python and C:** Provides Python-like syntax and dynamic typing while running in a minimal C footprint suitable for the RP2040.
- **Supports Systems Programming:** Unlike pure Python, SageLang includes explicit garbage collection control, class-based OOP, and low-level value representation suitable for embedded contexts.
- **Phases Incrementally:** The language grows through discrete, completed phases (1–18), each adding a cohesive feature set without breaking prior functionality.
- **Prioritizes Correctness:** Uses explicit memory management (via mark-and-sweep GC), scoped environments, and exception handling rather than relying on OS-level resource management.

1.2 Language Characteristics

Feature	Details
Syntax	Python-like: indentation-based blocks, <code>:</code> terminators, keywords like <code>let</code> , <code>var</code> , <code>proc</code> , <code>class</code>
Type System	Dynamically typed; values carry runtime type tags (numbers, strings, bools, arrays, dicts, tuples, classes, generators)
Scoping	Lexical scoping via nested environments. Note: Procedure calls matching VM built-in names resolve to module procedures first, causing recursion exceptions unless name shadowing is avoided.
Memory	Mark-and-sweep (tracing), ARC (reference counting), and ORC (optimized reference counting with cycle detection) modes
OOP	Class-based inheritance with single parent, methods, instance fields, <code>self</code> parameter, <code>super.init()</code> auto-self
Control Flow	<code>if/else</code> , <code>while</code> , <code>for...in</code> , <code>break</code> , <code>continue</code> , <code>return</code> , <code>try/catch/finally</code> , <code>raise</code> , <code>yield</code> , <code>defer</code>
Data Structures	Arrays (dynamic), dicts (string-keyed), tuples (immutable), slicing, indexing
Functions	First-class <code>proc</code> declarations with closures; anonymous <code>proc</code> expressions (<code>proc(x): body end</code>); native C functions
Generators	Full <code>yield</code> support with resumable state; <code>next()</code> function to iterate
Exceptions	<code>try/catch/finally/raise</code> with explicit exception objects and message strings

Feature	Details
Modules	<code>import module</code> , <code>import module as alias</code> , <code>from module import x</code> , <code>y</code> , <code>from module import x as y</code> . The same module can be imported under multiple bindings.
Standard Library	Native modules: <code>math</code> , <code>io</code> , <code>string</code> , <code>sys</code> , <code>thread</code> , <code>fat</code> , <code>socket</code> , <code>tcp</code> , <code>http</code> , <code>ssl</code>
Networking	POSIX sockets, TCP, HTTP/HTTPS (libcurl), SSL/TLS (OpenSSL)
Concurrency	Multi-threaded <code>proc</code> , <code>async/await</code> , <code>atomics</code> , <code>semaphores</code> , <code>condvars</code> , <code>rwlocks</code> , SMP detection
GPU Graphics	Vulkan + OpenGL 4.5 backends, handle-based API, 100+ functions, Android support
UI Widgets	Immediate-mode GUI: windows, panels, buttons, sliders, menus, text inputs
Compilation	C backend, LLVM IR (with GPU support), native assembly (x86-64, aarch64, rv64, mips), JIT self-extracting executables (with module bundling), plus initial profile suffixes for bare-metal / OSdev / UEFI targets

LLVM codegen has an additional import optimization path: both the C LLVM backend and the self-hosted LLVM backend resolve `from module import CONST` for foldable top-level constants during code generation (including `as` aliases).

1.3 Execution Model

SageLang uses a shared front-end with multiple execution backends:

1. **Source Code** → **Lexer** (tokenization with indentation tracking)
2. **Tokens** → **Parser** (recursive descent, builds AST)
3. **AST** → one of five backends:
 - **AST interpreter** (tree-walking, default)
 - **Bytecode compiler + VM** (stack-based, faster for hot loops)
 - **SGVM binary** (`--sgvm`, binary bytecode artifact)
 - **Lily Transpiler** (`--compile-to-lily`, `--compile-from-lily`)
 - **C codegen** (`--emit-c` / `--compile`)
 - **LLVM IR** (`--emit-llvm` / `--compile-llvm`, with GPU support)
 - **Native assembly** (`--emit-asm` / `--compile-native`, x86-64/aarch64/rv64/mips)
 - **Freestanding ELF** (`--compile-bare`, bare-metal kernel output)
 - **UEFI PE** (`--compile-uefi`, EFI application output)
 - **SageMetal VM** (`make metal-vm`, freestanding bytecode object)
4. **Runtime Execution**: Execution by the MetalVM engine integrated into `sage` or via standalone `sgvm` tool.
5. **Runtime Values** stored in the shared **heap** managed by **GC**

All execution modes share the same object model: a **global environment**, nested **child environments** for scopes, and **tagged Value objects** that are either immediate (numbers, bools) or GC-managed heap values (arrays, dicts, strings, classes, instances, functions, generators).

For LLVM-compiled workloads, `from module import CONST` is resolved at compile time for foldable top-level module constants (numbers/strings/bools/nil plus simple constant expressions). In the C LLVM backend, GPU module constants are also resolved at compile time and GPU calls emit direct bridges to the pure

C GPU API (`sgpu_*` in `gpu_api.h`), supporting both Vulkan and OpenGL backends. The bytecode VM provides 30 dedicated GPU opcodes for frame-loop hot paths.

The bytecode VM operates in hybrid mode by default: expressions, variables, loops (including break/continue), and function calls compile to stack bytecode, while unsupported constructs (classes, imports, exceptions, generators) fall back to the AST interpreter via `BC_OP_EXEC_AST_STMT`. This gives measurable speedups on loop-heavy workloads while maintaining full language coverage. Use `sage --runtime bytecode` or `sage --runtime auto` to enable.

1.4 Performance Characteristics

Sage ships with a benchmark suite (`benchmarks/01_fibonacci.sage` through `10_primes_sieve.sage`) with paired Python 3 implementations. Run `make benchmark-python` to compare. Typical results on the same workloads:

- **LLVM compiled:** 2-8x faster than CPython (fibonacci 3.9x, loop sum 7.7x)
- **C compiled:** 2-4x faster than CPython on most workloads
- **AST interpreter:** on par with CPython (faster on string/array ops, slower on recursion)
- **Bytecode VM:** similar to AST, faster on tight loops due to reduced dispatch overhead

Part 2: Internal Architecture and Core Modules

2.1 Module Dependency Graph

```
main.c
|- lexer.c / lexer.h           [Tokenization, indentation tracking]
|- parser.c / parser.h         [AST construction via recursive descent]
|- interpreter.c / interpreter.h [Tree-walking evaluation]
|- value.c / value.h           [Runtime value representation]
|- env.c / env.h               [Lexical scoping via linked-list environments]
|- gc.c / gc.h                 [Mark-and-sweep garbage collection]
|- ast.c / ast.h               [AST node factory functions]
|- token.h                     [Token type enumeration]
|- module.c / module.h         [Module loading, caching, imports]
|- compiler.c                   [C code generation backend]
|- llvm_backend.c               [LLVM IR generation (with GPU support)]
|- llvm_runtime.c               [Standalone runtime for LLVM-compiled programs]
|- codegen.c / codegen.h       [Native assembly (x86-64, aarch64, rv64, mips)]
|- graphics.c / graphics.h     [Vulkan GPU module for interpreter]
|- gpu_api.c / gpu_api.h       [Pure C GPU API (Vulkan + OpenGL)]
`- src/vm/                      [Bytecode VM: bytecode.c, vm.c, program.c, runtime.c]
```

2.2 Lexer (lexer.c / lexer.h)

Responsibility: Convert source text into a stream of **tokens** while tracking indentation levels (Python-style).

Key Data Structures: - **Token:** { TokenType type, const char* start, int length, int line }
- **Lexer State:** { const char* start, current, int line, int at_beginning_of_line } - **Indent Stack:** Array tracking nesting depth; generates TOKEN_INDENT / TOKEN_DEDEDENT automatically

Token Types (from token.h): - **Keywords:** and, as, async, await, break, case, catch, class, comptime, continue, default, defer, elif, else, end, enum, false, finally, for, from, if, import, in, init, let, macro, match, nil, not, or, print, proc, quote, raise, return, self, struct, super, trait, true, try, unquote, unsafe, var, while, yield, @ - *Note: `print`, `end`, `match`, `init`, `enum`, `struct`, and `trait` are soft keywords and can also be used as variable, property, or method names.* - **Operators:** +, -, *, /, =, ==, !=, <, >

>, <=, >=, and, or, &, |, ^, ~, <<, >> - *Note: Logical negation requires the **not** keyword. Using **!** produces a syntax error: Unexpected '!' (use 'not' for logical negation).* - **Punctuation:** (,), [,], {, }, :, ,, . - **Structural:** INDENT, DEDENT, NEWLINE, DOC_COMMENT, EOF, ERROR

Indentation Handling (crucial for Python-like syntax):

Resource Management: - **SAGE_MAX_READ_SIZE:** A 100MB limit is enforced on all file I/O and network receive operations to prevent memory exhaustion DoS attacks.

- Tracks column position at the start of each line.
- Compares current indent to stack top:
 - **Increase:** Push new level, emit `TOKEN_INDENT`.
 - **Decrease:** Pop levels, emit `TOKEN_DEDENT` for each popped level.
 - **Mismatch:** Error if indent doesn't align with a previous level.

Example:

```
proc greet(name):
    let msg = "Hello, " + name
    print msg
# Tokens: proc, identifier(greet), (, identifier(name), ), :, newline,
#         indent, let, identifier(msg), =, string, newline,
#         print, identifier(msg), newline,
#         dedent, eof
```

2.3 Parser (parser.c)

Responsibility: Build an **Abstract Syntax Tree (AST)** from tokens using **recursive descent parsing**.

Parsing Levels (operator precedence):

expression	→ assignment
assignment	→ logical_or ("=" assignment)?
logical_or	→ logical_and ("or" logical_and)*
logical_and	→ equality ("and" equality)*
equality	→ comparison (("==" "!=") comparison)*
comparison	→ addition (<"<" ">" "<=" ">=") addition)*
addition	→ term (<"+" "-") term)*
term	→ unary (<"*" "/"> unary)*
unary	→ (<"-" unary> postfix
postfix	→ primary (<"[" ... "]" "." property>)*
primary	→ NUMBER STRING BOOL NIL "<(" expr ">)" "<[" array ">)" "<{" dict ">)" IDENTIFI

Statement Parsing (top-down recursive descent):

statement	→ if_stmt while_stmt for_stmt return_stmt break continue try_stmt raise_stmt
declaration	→ class_decl proc_decl let_decl statement

Key AST Node Types (see `ast.h/ast.c`): - **Expr** variants: `EXPR_NUMBER`, `EXPR_STRING`, `EXPR_BOOL`, `EXPR_NIL`, `EXPR_VARIABLE`, `EXPR_CALL`, `EXPR_BINARY`, `EXPR_UNARY`, `EXPR_INDEX`, `EXPR_SLICE`, `EXPR_ARRAY`, `EXPR_DICT`, `EXPR_TUPLE`, `EXPR_GET`, `EXPR_SET` - **Stmt** variants: `STMT_EXPRESSION`, `STMT_LET`, `STMT_PRINT`, `STMT_IF`, `STMT_WHILE`, `STMT_FOR`, `STMT_BLOCK`, `STMT_PROC`, `STMT_CLASS`, `STMT_RETURN`, `STMT_BREAK`, `STMT_CONTINUE`, `STMT_TRY`, `STMT_RAISE`, `STMT_YIELD`, `STMT_DEFER`, `STMT_MATCH`, `STMT_IMPORT`

Example Parsing:

```
Input: let x = 10 + 5
Output: Stmt {
    type: STMT_LET,
    name: Token("x"),
```

```

    initializer: Expr {
      type: EXPR_BINARY,
      left: Expr { type: EXPR_NUMBER, value: 10 },
      op: Token(PLUS),
      right: Expr { type: EXPR_NUMBER, value: 5 }
    }
  }
}

```

2.4 Value System (value.h / value.c)

Responsibility: Define and manage **runtime value types** and their operations.

Core Value Type:

```

struct Value {
  ValueType type; // VAL_NUMBER, VAL_STRING, VAL_ARRAY, etc.
  union {
    double number;
    int boolean;
    char* string;
    NativeFn native;
    FunctionValue* function;
    ArrayValue* array;
    DictValue* dict;
    TupleValue* tuple;
    ClassValue* class_val;
    InstanceValue* instance;
    ExceptionValue* exception;
    GeneratorValue* generator;
  } as;
};

```

Value Types:

Type	Storage	Mutable	Use Case
VAL_NUMBER	double (on stack)	Yes	Integers, floats; 10, 3.14
VAL_BOOL	int (on stack)	Yes	Boolean logic; true, false
VAL_NIL	N/A	No	Null value; nil
VAL_STRING	char* (heap)	No	Text; immutable (assign new to mutate)
VAL_ARRAY	ArrayValue* (dynamic)	Yes	Lists; [1, 2, 3], push(), indexing
VAL_DICT	DictValue* (string-keyed)	Yes	Maps; {"key": value}, dict_set()
VAL_TUPLE	TupleValue* (immutable)	No	Fixed-size sequences; (1, 2, 3)
VAL_FUNCTION	FunctionValue* (closure)	No	User-defined procs with captured environment
VAL_NATIVE	NativeFn (C function ptr)	No	Built-in functions like len(), print()
VAL_CLASS	ClassValue* (metadata)	No	Class definition with methods and parent
VAL_INSTANCE	InstanceValue* (fields dict)	Yes	Object; instance of a class

Type	Storage	Mutable	Use Case
VAL_EXCEPTION	ExceptionValue* (message)	No	Exception; raised via <code>raise</code>
VAL_GENERATOR (SAGE_TAG_GENERATOR)	GeneratorValue* (Resumable)	Yes	Iterable; maintains state across <code>yield</code> , supports native yield collector routines and <code>next()</code> builtin dispatch

Array Operations: - `array_push(arr, val)`: Append to dynamic array (resizes if needed). - `array_get(arr, index)`, `array_set(arr, index, val)`: Access/modify. - `array_slice(arr, start, end)`: Subarray (negative indices supported).

Dictionary Operations ($O(1)$ amortized via hash table): - `dict_set(dict, key, value)`: Insert/update using FNV-1a hashing with open-addressing and linear probing. Table grows at 75% load factor. - `dict_get(dict, key)`, `dict_has(dict, key)`: $O(1)$ average lookup. - `dict_delete(dict, key)`: Backward-shift deletion preserves probe chain integrity (no tombstones). - `dict_keys(dict)`, `dict_values(dict)`: Return as arrays (iterates capacity, skips empty slots).

String Operations: - `string_split(str, delim)`, `string_join(arr, sep)`: Splitting/joining arrays. - `string_replace(str, old, new)`, `string_upper/lower/strip(str)`: Transformation.

Class/Instance Operations: - `class_create(name, parent)`: Define a class with optional parent. - `class_add_method(class, name, method_stmt)`: Add methods. - `instance_create(class_def)`: Create instance with empty field dictionary. - `instance_set_field()`, `instance_get_field()`: Access instance variables.

2.5 Environment and Scoping (env.h / env.c)

Responsibility: Implement **lexical scoping** via a chain of nested environments (linked list).

Environment Structure:

```

struct Env {
    EnvNode* head;           // Variables defined in this scope (linked list)
    struct Env* parent;      // Enclosing scope
    struct Env* alloc_next;  // Internal registry for GC/shutdown cleanup
    int marked;              // GC mark flag (0 = unmarked, 1 = reachable)
};

struct EnvNode {
    char* name;              // Variable name
    Value value;             // Stored value
    struct EnvNode* next;
};

```

Scoping Rules: - `env_define(env, name, length, value)`: Create or update a variable **in the current scope only**. If variable exists in current scope, update it; otherwise, create new. - `env_get(env, name, length, value*)`: Look up a variable, searching current scope then parent chains recursively. Return 1 if found, 0 if not. - `env_assign(env, name, length, value)`: Assign to an **existing** variable by searching up the scope chain. Used for reassigning already-defined variables.

Example Scoping:

```

Global: { x: 10, y: 20 }
  Block: { z: 30 } → parent: Global
    Inner: { x: 100 } → parent: Block

```

In Inner scope:

- `env_get("x")` → 100 (found in current scope)
- `env_get("z")` → 30 (found in parent Block)
- `env_get("y")` → 20 (found in Global via parent chain)

2.6 Garbage Collector (gc.h / gc.c)

Responsibility: Automatically reclaim heap memory using tracing or reference counting.

GC Modes:

- **Tracing** (`--gc:tracing`, default): Concurrent tri-color mark-sweep with SATB write barriers. Best for general use with sub-millisecond pauses.
- **ARC** (`--gc:arc`): Deterministic reference counting. High performance for linear object graphs; does not collect cycles.
- **ORC** (`--gc:orc`): Optimized Reference Counting with Lins' trial deletion cycle collector. Combines the determinism of ARC with a robust cycle detection algorithm for complex graphs.

Tracing GC Overview: - **Mark Phase:** Starting from **root set** (global environment, function registry, call stack), mark all reachable objects. Born-black objects during concurrent marking. - **Sweep Phase:** Free all unmarked objects and add them to a free pool. - **Triggering:** Automatic based on object count and heap pressure, or manual via `gc_collect()`.

GC Configuration:

```
#define GC_HEAP_SIZE (1024 * 1024) // 1 MB theoretical heap
#define GC_THRESHOLD 1000         // Collect after 1000 object allocations
```

GC Header (prepended to all heap allocations):

```
struct GCHeader {
    int marked;    // 1 = marked (reachable), 0 = unmarked (candidate for freeing)
    int type;      // Object type (VAL_ARRAY, VAL_DICT, etc.)
    void* next;    // Linked list of all allocated objects
};
```

Root Set: - Global environment (`g_global_env`) and its entire scope chain. - Function registry (linked list of all defined proc declarations). - Call stack (environments of active function calls).

GC Statistics (via `gc_stats()` native function):

```
bytes_allocated: Total heap allocated (cumulative)
num_objects: Current number of live objects
collections: Number of GC runs so far
objects_freed: Objects freed in most recent collection
next_gc: Objects until next collection threshold
```

Usage:

```
# Automatic collection when threshold hit
let arr = [1, 2, 3] # Allocates array

# Manual trigger
gc_collect()

# Check stats
let stats = gc_stats()
print stats
# Output: {"bytes_allocated": 1024, "num_objects": 5, "collections": 2, ...}
```

```
# Control GC
gc_disable() # Pause automatic collection
gc_enable()  # Resume automatic collection
```

2.7 Interpreter (interpreter.c / interpreter.h)

Responsibility: Tree-walking evaluation of the AST, maintaining runtime state and executing statements/expressions.

Execution Result (tracks control flow):

```
struct ExecResult {
    Value value;           // Return value
    int is_returning;      // `return` statement hit
    int is_breaking;       // `break` statement hit
    int is_continuing;     // `continue` statement hit
    int is_throwing;       // Exception raised
    Value exception_value; // Exception object
    int is_yielding;       // `yield` hit (generators)
    void* next_stmt;       // Resume point for generators
};
```

Statement Execution (`interpret(stmt, env)`): - **STMT_LET**: Define variable (`let` or `var`) in current scope; evaluate initializer if present. - **STMT_EXPRESSION**: Evaluate expression for side effects. - **STMT_PRINT**: Evaluate expression and print value. - **STMT_IF**: Evaluate condition; execute then-branch or else-branch. - **STMT_WHILE**: Loop while condition is truthy; respect `break/continue`. - **STMT_FOR**: Iterate over array/range; update loop variable in scope. - **STMT_BLOCK**: Create child scope, execute statements in sequence. - **STMT_PROC**: Register function in function registry (not executed). - **STMT_CLASS**: Create class object; register methods. - **STMT_RETURN**: Exit function with value. - **STMT_TRY**: Execute try-block; on exception, search matching catch-clause; always run finally-block. - **STMT_RAISE**: Throw exception (string or exception object). - **STMT_YIELD**: In generators, pause execution and return value; next `next()` call resumes. - **STMT_IMPORT**: Load module; populate current environment with imports.

Expression Evaluation (`eval_expr(expr, env)`): - **Literals**: Numbers, strings, bools, `nil` returned as-is. - **Variables**: Look up in environment via `env_get()`. - **Binary Ops**: Evaluate left, then right (short-circuit for `and/or` in both interpreter and C backend); apply operator. - **Arithmetic**: `+`, `-`, `*`, `/` on numbers; `+` on strings (concatenation). - **Comparison**: `<`, `>`, `<=`, `>=` on numbers; `==`, `!=` on any type. - **Logical**: `and`, `or` with short-circuit evaluation (in both interpreter and C backend). - **Calls**: Look up function (user-defined or native); bind arguments to parameters; execute in new environment; return result. - **Indexing**: Evaluate array and index; return element or slice if range. - **Array/Dict/Tuple Construction**: Evaluate elements; construct heap-allocated structure. - **Property Access**: For instances, look up field in instance's field dict.

Function Call Mechanics: 1. Resolve callee to `ProcStmt*` or `NativeFn`. 2. Create **child environment** with current scope as parent. 3. Bind parameters to arguments via `env_define()`. 4. Execute function body statements. 5. On `return`, break with value; otherwise use `nil`. 6. Return to caller with result.

Generator Execution (Phase 7): 1. `proc my_gen()`: creates a function that yields values. 2. Call `my_gen()` → returns `VAL_GENERATOR` value. 3. Call `next(gen)` → initializes generator environment, runs until first `yield`. 4. Each `next()` resumes from saved `current_stmt`; yields next value or `nil` if exhausted.

Native Function Registry: SageLang provides built-in functions injected into global environment via `init_stdlib(env)`:

Function	Signature	Purpose
<code>len(arr/str/dict)</code>	<code>Value → number</code>	Get length of collection
<code>push(arr, val)</code>	<code>(array, value) → nil</code>	Append to array

Function	Signature	Purpose
pop(arr)	array → value	Remove and return last element
range(start, end)	(number, number) → array	Generate list [start, start+1, ..., end-1]
slice(arr/str, start, end)	(array/string, number, number) → array/string	Extract subarray or substring
split(str, delim)	(string, string) → array	Split string by delimiter
join(arr, sep)	(array, string) → string	Join array elements with separator
replace(str, old, new)	(string, string, string) → string	Replace all occurrences
upper(str)	string → string	Uppercase
lower(str)	string → string	Lowercase
strip(str)	string → string	Trim whitespace
str(val)	value → string	Convert to string
tonumber(str)	string → number	Parse number
input()	() → string	Read line from stdin
clock()	() → number	Elapsed seconds since process start
dict_keys(dict)	dict → array	Get keys as array
dict_values(dict)	dict → array	Get values as array
dict_has(dict, key)	(dict, string) → bool	Check if key exists
dict_delete(dict, key)	(dict, string) → nil	Remove key from dict
hasattr(obj, name)	(instance dict, string) → bool	Duck-typing check for instance methods, fields, and dictionary keys
val_tag(val)	value → number	Get internal numeric tag identifier of a value
build_info()	() → dict	Get binary build metadata (version, arch, build type, spec version)
type(val)	value → string	Get string name of type (type(val) returns “bytes” for byte buffers)
chr(n)	number → string	Get single-character string from ASCII code
ord(c)	string → number	Get ASCII code of single-character string
hash(val)	value → number	Get hash code of a value
sizeof(val)	value → number	Get memory size of a value in bytes
doc(val)	value → string	Retrieve documentation from a function
gc_collect()	() → nil	Trigger garbage collection
gc_stats()	() → dict	Get GC statistics
gc_enable(), gc_disable()	() → nil	Control automatic GC
next(gen)	generator → value	Resume generator, get next yielded value
ffi_open(path)	string → clib	Open shared library via dlopen
ffi_call(lib, func, ret, ...)	(clib, string, string, ...) → value	Call C function in library

Function	Signature	Purpose
<code>ffi_sym(lib, name)</code>	<code>(clib, string) → bool</code>	Check if symbol exists in library
<code>ffi_close(lib)</code>	<code>clib → nil</code>	Close shared library handle
<code>mem_alloc(size)</code>	<code>number → pointer</code>	Allocate raw memory (zero-initialized)
<code>mem_free(ptr)</code>	<code>pointer → nil</code>	Free allocated memory
<code>mem_read(ptr, off, type)</code>	<code>(pointer, number, string) → value</code>	Read value at ptr+offset
<code>mem_write(ptr, off, type, val)</code>	<code>(pointer, number, string, value) → nil</code>	Write value at ptr+offset
<code>mem_size(ptr)</code>	<code>pointer → number</code>	Get allocation size
<code>addressof(val)</code>	<code>value → number</code>	Get memory address of a value
<code>asm_exec(code, ret, ...)</code>	<code>(string, string, ...) → value</code>	Compile and execute assembly
<code>asm_compile(code, arch, out)</code>	<code>(string, string, string) → bool</code>	Cross-compile assembly to object file
<code>asm_arch()</code>	<code>() → string</code>	Get host architecture name
<code>struct_def(fields)</code>	<code>array → dict</code>	Define C struct layout with alignment
<code>struct_new(def)</code>	<code>dict → pointer</code>	Allocate zeroed struct instance
<code>struct_get(ptr, def, name)</code>	<code>(pointer, dict, string) → value</code>	Read struct field
<code>struct_set(ptr, def, name, val)</code>	<code>(pointer, dict, string, value) → nil</code>	Write struct field
<code>struct_size(def)</code>	<code>dict → number</code>	Get total struct size
<code>cpu_count()</code>	<code>() → number</code>	Logical CPU core count
<code>cpu_physical_cores()</code>	<code>() → number</code>	Physical CPU core count
<code>cpu_has_hyperthreading()</code>	<code>() → bool</code>	HT check
<code>thread_set_affinity(id)</code>	<code>number → number</code>	Pin current thread to core
<code>thread_get_core()</code>	<code>() → number</code>	Current CPU core ID
<code>atomic_new(val)</code>	<code>value → pointer</code>	Create atomic long
<code>atomic_load(a)</code>	<code>pointer → number</code>	Atomic read
<code>atomic_store(a, v)</code>	<code>(pointer, number) → nil</code>	Atomic write
<code>atomic_add(a, v)</code>	<code>(pointer, number) → number</code>	Atomic fetch-and-add
<code>atomic_cas(a, e, d)</code>	<code>(pointer, num, num) → bool</code>	Atomic compare-and-swap
<code>atomic_exchange(a, v)</code>	<code>(pointer, number) → number</code>	Atomic fetch-and-exchange
<code>sem_new(n)</code>	<code>number → pointer</code>	Create POSIX semaphore
<code>sem_wait(s)</code>	<code>pointer → nil</code>	Blocking wait
<code>sem_post(s)</code>	<code>pointer → nil</code>	Signal/release
<code>sem_trywait(s)</code>	<code>pointer → bool</code>	Non-blocking wait

2.8 Module System (module.h / module.c)

Responsibility: Load `.sage` files as reusable modules and manage imports.

Module Structure:

```

struct Module {
    char* name;           // Module identifier (e.g., "math")
    char* path;           // Full filesystem path to .sage file
    Environment* env;      // Module's own environment (exports)
    bool is_loaded;        // Execution complete?
    bool is_loading;       // Currently executing? (cycle detection)

```

```

    struct Module* next;          // For linked-list cache
};

struct ModuleCache {
    Module* modules;              // All loaded modules
    char** search_paths;          // Directories to search (., ./lib, ./modules)
    int search_path_count;
};

```

Module Lifecycle: 1. `load_module(cache, "math")`: Search in `search_paths` for `math.sage` or `math/__init__.sage`; create `Module` struct; add to cache. 2. `execute_module(module, global_env)`: Read file; lex/parse/interpret in module's own environment (child of `global`). 3. **Exports**: Any top-level `proc`, `class`, `let` in module becomes available for import.

Import Statements (parsed in `parser.c`, executed in `interpreter.c`):

```

# Import entire module (unaliased)
import math
# → Available as: math (currently placeholder; future: module.sin(), etc.)

# Import module with alias
import math as m
# → Available as: m

# Import specific items
from math import sin, cos
# → sin and cos directly available in scope

# Import specific items with aliases
from math import sin as sine, cos as cosine
# → sine, cosine available in scope

# Star import (all exports)
from math import *
# → All module exports available directly

```

Search Path Resolution: 1. Try `{search_path}/module_name.sage`. 2. Try `{search_path}/module_name/__init__.sage`. 3. If not found in any search path, error.

Default Search Paths:

```

- "."          (current directory)
- "./lib"      (local lib subdirectory)
- "./modules"  (local modules subdirectory)

```

Import Processing: - `import_all(env, "math")`: Load module, define module name in `env`. - `import_from(env, "math", items[], count)`: Load module, look up items, define in `env`. - `import_as(env, "math", "m")`: Load module, define alias in `env`.

Compiled LLVM Constant Imports: - Applies to both LLVM pipelines: the C backend (`src/c/llvm_backend.c`) and the self-hosted backend (`src/sage/llvm_backend.sage`). - `from module import NAME` (and `from module import NAME as ALIAS`) is pre-resolved from module top-level `let` constants when they are compile-time foldable. - Supported imported value shapes are scalar constants (`number`, `string`, `bool`, `nil`) and expressions composed from previously resolved constants. - Search paths for this compile-time resolution match module lookup defaults: `./`, `./lib`, and `./modules`. - In the C LLVM backend, `from gpu import CONST` is also resolved from the built-in GPU constant table. - Unresolved imported constants are treated as compile errors in LLVM codegen rather than generating unresolved `%NAME` loads.

Circular Dependency Detection: `is_loading` flag prevents infinite loops during module execution.

Part 3: Language Features by Phase

Phase 1–3: Core Language (Arithmetic, Variables, Control Flow)

Implemented: Basic arithmetic, variable declaration, conditionals, loops, functions.

Example:

```
let x = 10
let y = 20
print x + y # Output: 30

if x > 5:
    print "x is greater than 5"
else:
    print "x is not greater than 5"

proc add(a, b):
    return a + b

print add(3, 4) # Output: 7
```

Phase 4: Garbage Collection

Added: Mark-and-sweep GC automatically manages heap.

Example:

```
let arr = [1, 2, 3, 4, 5]
print len(arr) # 5

# GC runs automatically when threshold hit
gc_collect() # Force collection
let stats = gc_stats()
print stats["collections"] # Number of GC runs
```

Phase 5–6: Data Structures and OOP

Phase 5: Arrays, dicts, tuples, slicing, indexing.

```
let arr = [1, 2, 3, 4, 5]
print arr[0] # 1
print arr[1:3] # [2, 3]
arr[0] = 10
print arr # [10, 2, 3, 4, 5]

let dict = {"name": "Alice", "age": 30}
print dict["name"] # "Alice"
dict["age"] = 31

let tuple = (1, 2, 3)
print tuple # (1, 2, 3) # Immutable
```

Phase 6: Classes, inheritance, instance methods.

```

class Animal:
  proc init(name):
    self.name = name
    self.age = 0

  proc birthday():
    self.age = self.age + 1

  proc speak():
    print self.name + " makes a sound"

class Dog(Animal):
  proc speak():
    print self.name + " barks"

let dog = Dog("Buddy")
print dog.name          # "Buddy"
dog.birthday()
print dog.age           # 1
dog.speak()             # "Buddy barks"

```

Phase 7: Advanced Control Flow, Exceptions, Generators

Exceptions:

```

proc divide(a, b):
  if b == 0:
    raise "Division by zero"
  return a / b

try:
  print divide(10, 0)
catch e:
  print "Caught error: " + e
finally:
  print "Cleanup done"

```

Generators:

```

proc count_up(n):
  let i = 0
  while i < n:
    yield i
    i = i + 1

let gen = count_up(3)
print next(gen)      # 0
print next(gen)      # 1
print next(gen)      # 2
print next(gen)      # nil (exhausted)

```

Deferred Code (Phase 7 addition): `defer` executes a statement on scope exit, in LIFO (Last-In, First-Out) order, which is extremely useful for cleanup:

```

proc read_and_print():
  let f = file_open("data.txt")

```

```

defer file_close(f)
print file_read(f)

```

Pattern Matching (Phase 7 addition): `match` allows matching a value against multiple patterns, with optional guards and a default fallback:

```

let x = 42
match x:
  case 1:
    print "one"
  case 42 if y > 0:
    print "forty-two with positive y"
  case 42:
    print "forty-two"
  default:
    print "other"

```

Phase 8: Modules

```

# math.sage
proc sin(x):
  # Approximate sine
  return x # Stub

proc cos(x):
  return 1 # Stub

# main.sage
import math as m
print m.sin(0) # Calls math.sin in module's namespace

# Or:
from math import sin, cos
print sin(0)
print cos(0)

```

Phase 9: Low-Level Programming

Bitwise Operators:

```

# Bitwise AND, OR, XOR
print 5 & 3      # 1  (0101 & 0011 = 0001)
print 5 | 3      # 7  (0101 | 0011 = 0111)
print 5 ^ 3      # 6  (0101 ^ 0011 = 0110)

# Bitwise NOT
print ~0         # -1 (all bits flipped)

# Shift operators
print 1 << 4      # 16 (shift left by 4)
print 16 >> 2     # 4  (shift right by 2)

# Practical: extract lower nibble
let val = 255
let mask = 15
print val & mask  # 15

```

```
# Practical: check if bit is set
let x = 7
let bit2 = (x >> 2) & 1
print bit2 == 1    # true
```

Bitwise Safety: Shift amounts are validated at runtime — values outside 0-63 return 0 instead of causing C undefined behavior. Right-shift is arithmetic (sign-extending) on signed values. Floating-point operands are truncated to `long long` before bitwise operations.

Unsafe Blocks: As of version 4.0.1 (Specification 2.0), low-level memory operations or operations bypassing typical safety checks should be wrapped in an `unsafe` block. `unsafe` blocks must be terminated with the `end` keyword.

```
unsafe:
  # Memory primitives or FFI calls here
  let ptr = mem_alloc(16)
  mem_write(ptr, 0, "int", 42)
  mem_free(ptr)
end
```

Metaprogramming (Phase 17): SageLang includes experimental metaprogramming syntax constructs evaluated during parsing and compilation: - **`comptime(expr)`**: Evaluates the inner expression at compile-time (constant folding). - **`comptime::`**: A block statement executing code directly within the compiler context. - **`macro name(args)::`**: Declares an AST transformation macro. - **`quote` / `unquote`**: Reserved lexer keywords for future AST token quoting integration.

Note: Phase 17 features are implemented in the parser but have incomplete compiler backend propagation outside constant expressions.

Foreign Function Interface (FFI):

```
# Open a shared C library
let libm = ffi_open("libm.so.6")

# Call C functions with return type and arguments
let result = ffi_call(libm, "sqrt", "double", 144.0)
print result    # 12

let s = ffi_call(libm, "sin", "double", 0.0)
print s        # 0

# Check if a symbol exists
print ffi_sym(libm, "cos")    # true
print ffi_sym(libm, "bogus")  # false

# Close when done
ffi_close(libm)

# Call libc functions
let libc = ffi_open("libc.so.6")
let n = ffi_call(libc, "strlen", "long", "hello")
print n    # 5
ffi_close(libc)
```

FFI supports return types: `"double"`, `"int"`, `"long"`, `"string"`, `"void"`, with up to 3 arguments (numeric or string). Passing more than 3 arguments returns an error. Library handles are tracked by the GC and properly freed on collection.

Raw Memory Operations:

```
# Allocate 32 bytes of raw memory
let buf = mem_alloc(32)
print mem_size(buf)      # 32

# Write and read different types
mem_write(buf, 0, "byte", 42)
print mem_read(buf, 0, "byte")    # 42

mem_write(buf, 4, "int", 12345)
print mem_read(buf, 4, "int")     # 12345

mem_write(buf, 8, "double", 3.14)
print mem_read(buf, 8, "double")  # 3.14

# Get memory address of any value
let arr = [1, 2, 3]
print addressof(arr)    # memory address as number

# Free when done
mem_free(buf)
```

The `mem_read(ptr, offset, type)` function requires exactly 3 arguments, and `mem_write(ptr, offset, type, value)` requires exactly 4 arguments. Supported explicit type strings for `mem_read/mem_write` are: "byte" (1 byte), "int" (4 bytes), "double" (8 bytes), "string" (read-only, null-terminated). Allocations are capped at 64MB. Negative offsets are rejected. Bounds checking is enforced for owned memory (offset + type size must not exceed allocation). Double-free is prevented via handle nullification. Memory pointers are GC-tracked and freed on collection if owned.

Inline Assembly (x86-64, aarch64, rv64, mips):

```
# Detect host architecture
print asm_arch()      # "x86_64"

# Execute x86-64 assembly: return a constant
let val = asm_exec("    mov $42, %rax", "int")
print val              # 42

# Add two integers (args in rdi, rsi per System V ABI)
let sum = asm_exec("    mov %rdi, %rax\n    add %rsi, %rax", "int", 10, 32)
print sum              # 42

# Add two doubles (args in xmm0, xmm1)
let r = asm_exec("    addsd %xmm1, %xmm0", "double", 1.5, 2.7)
print r                # 4.2

# Cross-compile for aarch64 (requires aarch64-linux-gnu-as)
let ok = asm_compile("    mov x0, #42", "aarch64", "/tmp/out.o")

# Cross-compile for RISC-V 64 (requires riscv64-linux-gnu-as)
let ok2 = asm_compile("    li a0, 42", "rv64", "/tmp/out_rv.o")
```

Supported architectures: "x86_64", "aarch64", "rv64", "mips". Return types: "int", "double", "void". Up to 4 numeric arguments.

C Struct Interop:

```

# Define a C-compatible struct: { int x; int y; double z; }
let Point = struct_def(["x", "int"], ["y", "int"], ["z", "double"])
print struct_size(Point)    # 16 (with alignment)

# Allocate and populate
let p = struct_new(Point)
struct_set(p, Point, "x", 10)
struct_set(p, Point, "y", 20)
struct_set(p, Point, "z", 3.14)

# Read fields
print struct_get(p, Point, "x")    # 10
print struct_get(p, Point, "y")    # 20
print struct_get(p, Point, "z")    # 3.14

mem_free(p)

# Alignment example: { char a; double b; int c; }
# Layout: a@0, pad(7), b@8, c@16, pad(4) = 24 bytes
let S = struct_def(["a", "char"], ["b", "double"], ["c", "int"])
print struct_size(S)              # 24

Supported types: "char", "byte" (1), "short" (2), "int" (4), "long" (8), "float" (4), "double" (8),
"ptr" (8). Alignment follows C ABI rules.

```

Part 4: Writing SageLang Programs

4.1 Basic Syntax Patterns

Variables and Types:

let x = 42	# Variable binding
var y = 100	# Explicitly mutable binding
let s = "Hello, World"	# String
let b = true	# Boolean
let arr = [1, 2, 3]	# Array
let dict = {"a": 1, "b": 2}	# Dictionary
let tup = (1, "two", 3)	# Tuple (immutable)
let n = nil	# Nil (null)

Operators:

# Arithmetic	
print 10 + 5	# 15
print 10 - 5	# 5
print 10 * 5	# 50
print 10 / 2	# 5
print -5	# -5
# Comparison	
print 5 == 5	# true
print 5 != 3	# true
print 5 > 3	# true
print 5 < 10	# true

```
# Logical
print true and false    # false
print true or false     # true
```

```
# Bitwise
print 5 & 3              # 1
print 5 | 3              # 7
print 5 ^ 3              # 6
print ~0                 # -1
print 1 << 4             # 16
print 16 >> 2            # 4
```

String Operations:

```
let a = "Hello"
let b = "World"
print a + " " + b        # "Hello World"

let words = "a,b,c".split(",")
print words               # ["a", "b", "c"]

let joined = ["x", "y"].join("-")
print joined              # "x-y"

print "hello".upper()     # "HELLO"
print "HELLO".lower()     # "hello"
```

Arrays:

```
let arr = [10, 20, 30, 40, 50]
print len(arr)            # 5
print arr[0]              # 10
arr[0] = 100
print arr[0]              # 100
push(arr, 60)
print len(arr)            # 6

# Slicing
print arr[1:3]            # [20, 30]
print arr[2:]             # [30, 40, 50, 60] (from index 2 to end)
print arr[:2]             # [100, 20] (from start to index 2)
```

Dictionaries:

```
let dict = {"name": "Alice", "age": 30}
print dict["name"]        # "Alice"
dict["age"] = 31
print dict_has(dict, "name") # true
print dict_keys(dict)     # ["name", "age"]
print dict_values(dict)    # ["Alice", 31]

dict_delete(dict, "age")
print dict_has(dict, "age") # false
```

4.2 Control Flow Patterns

Conditionals:

```

let x = 15

if x < 10:
    print "small"
else:
    print "large"

# Nested conditions
if x > 10:
    if x > 20:
        print "very large"
    else:
        print "medium"
else:
    print "small"

Loops:

# While loop
let i = 0
while i < 5:
    print i
    i = i + 1

# For loop over array
let arr = [10, 20, 30]
for item in arr:
    print item

# For loop over range
for i in range(0, 5):
    print i

# Note: Strings are not currently iterable via `for` loops.
# Use a `while` loop with index-based access instead.
let s = "hello"
let i_str = 0
while i_str < len(s):
    print s[i_str]
    i_str = i_str + 1

# Break and continue
for i in range(0, 10):
    if i == 3:
        continue
    if i == 7:
        break
    print i

```

4.3 Functions and Closures

Function Definition and Call:

```

proc greet(name):
    print "Hello, " + name

```

```

greet("Alice")

proc add(a, b):
  return a + b

let result = add(5, 3)
print result          # 8

Closures (captured environment):
proc make_multiplier(factor):
  proc multiply(x):
    return x * factor  # Captures 'factor' from enclosing scope
  return multiply

let times3 = make_multiplier(3)
let times5 = make_multiplier(5)

print times3(10)      # 30
print times5(10)      # 50

```

4.4 Structs, Enums, Traits, and Classes

Struct Definition (Phase 1.7): SageLang provides native `struct` support as a lightweight alternative to classes. Structs automatically generate `init`, `to_string`, and equality methods based on their fields. Under the hood, structs are represented as standard class/instance values (`VAL_CLASS/VAL_INSTANCE`).

```

struct Point:
  x
  y
end

let p = Point(10, 20)
print p.x          # 10
print p            # Point(x=10, y=20)

```

Enum Definition (Phase 1.7): Native C-like `enum` support maps variant names to sequential integers starting from 0. At runtime, enums are represented as dictionaries. Note that tagged unions (ADTs with associated data) are still implemented via the standard library (`std.enum`).

```

enum Color:
  Red
  Green
  Blue

print Color.Red    # 0
print Color.Green  # 1

```

Trait Definition (Phase 1.7): Traits define an interface contract of method signatures. A trait compiles down to a dictionary containing the list of required method names.

```

trait Printable:
  proc to_string(self): pass
end

```

Class Definition:

```

class Point:
  proc init(x = 0, y = 0): # Default parameters are supported

```

```

        self.x = x
        self.y = y

    proc distance_from_origin():
        # Simplified (no sqrt)
        return self.x * self.x + self.y * self.y

let p = Point(3, 4)
print p.x           # 3
print p.y           # 4
print p.distance_from_origin() # 25

```

Inheritance:

```

class Vehicle:
    proc init(wheels):
        self.wheels = wheels

    proc describe():
        print "Vehicle with " + str(self.wheels) + " wheels"

class Car(Vehicle):
    proc init(wheels, doors):
        self.wheels = wheels
        self.doors = doors

    proc describe():
        print "Car with " + str(self.wheels) + " wheels and " + str(self.doors) + " doors"

let car = Car(4, 4)
car.describe()           # "Car with 4 wheels and 4 doors"

```

Calling Parent Methods with super:

Use `super.init(args)` to call the parent class constructor, and `super.method(args)` to call any parent method. `self` is automatically injected. This works with chained inheritance (3+ levels, e.g., $A \rightarrow B \rightarrow C$).

```

class Animal:
    proc init(self, name):
        self.name = name
    proc speak(self):
        print self.name + " speaks"

class Dog(Animal):
    proc init(self, name, breed):
        super.init(name)
        self.breed = breed
    proc speak(self):
        super.speak()
        print self.name + " barks"

let d = Dog("Rex", "Labrador")
d.speak()
# Rex speaks
# Rex barks

```

The `->` arrow operator can also be used with `super`: `super->init(args)`.

Arrow Operator (->):

The -> operator is a syntactic alias for . (dot), providing systems-language style pointer/member access. It works identically for field reads, field writes, and method calls.

```
let p = Point(3, 7)
print p->x           # same as p.x
p->x = 10             # same as p.x = 10
print p->to_string()  # same as p.to_string()
```

-> is interchangeable with . everywhere: attribute access, method calls, and **super** calls (**super->init(self, args)**).

4.5 Exception Handling

Try-Catch-Finally:

```
proc safe_divide(a, b):
  if b == 0:
    raise "Zero denominator"
  return a / b

try:
  let result = safe_divide(10, 0)
catch error:
  print "Error: " + error
finally:
  print "Done"
```

```
# Output:
# Error: Zero denominator
# Done
```

Finally Block Semantics: The finally block always executes, and its control flow takes precedence over try/catch. If finally contains **return**, **break**, **continue**, or **raise**, that overrides the try/catch result. If finally executes normally, the try/catch result is preserved. This matches Python/Java behavior.

Raise: You can raise any value. Strings become exception messages directly. Numbers, booleans, and nil are converted to their string representation. Non-string/non-exception values become “Unknown error”.

```
raise "file not found"      # Exception with message "file not found"
raise 404                   # Exception with message "404"
raise nil                   # Exception with message "nil"
```

4.6 Generators

Simple Generator:

```
proc fibonacci(n):
  let a = 0
  let b = 1
  let count = 0
  while count < n:
    yield a
    let temp = a + b
    a = b
    b = temp
    count = count + 1
```

```

let gen = fibonacci(5)
print next(gen)           # 0
print next(gen)           # 1
print next(gen)           # 1
print next(gen)           # 2
print next(gen)           # 3
print next(gen)           # nil

```

4.7 Modules

math.sage:

```
let PI = 3.14159
```

```

proc square(x):
    return x * x

```

```

proc cube(x):
    return x * x * x

```

```

proc factorial(n):
    if n <= 1:
        return 1
    return n * factorial(n - 1)

```

main.sage:

```
import math as m
```

```

print m.PI                # 3.14159
print m.square(5)         # 25
print m.cube(3)           # 27
print m.factorial(5)      # 120

```

Or:

```

from math import PI, square, cube
print PI
print square(5)
print cube(3)

```

4.8 Defer Statements

Defer schedules cleanup code to run when the enclosing block exits, regardless of whether exit is via normal flow, **return**, **break**, **continue**, or exception. Multiple defers execute in LIFO (last-in, first-out) order.

```

proc process_file(path):
    let handle = open(path)
    defer:
        close(handle)
        print("file closed")
    # ... work with handle ...
    # close() runs automatically when block exits

```

```

proc multi_defer():
    defer:
        print("third")
    defer:

```

```

    print("second")
defer:
    print("first")
print("working")
# Output: working, first, second, third

```

Defer works with all exit paths: - Normal block completion - **return** from a function - **break/continue** from a loop - Exceptions (defers run before exception propagates)

Implementation: the interpreter collects deferred statements in a 1024-slot stack per block. On block exit, they execute in reverse order.

4.9 Pattern Matching (match/case/default)

Match evaluates a value and compares it against case patterns. Patterns support literal value comparisons, bare-identifier variable bindings (which capture the matched value into a clause-scoped variable visible in guards and clause bodies), match guards (**if condition**), and the **_** wildcard pattern. An optional **default** clause runs if no case matches.

```

match x:
    case n if n > 10:
        print "Large number: " + str(n)
    case n:
        print "Matched value: " + str(n)
    default:
        print "Fallback"

let cmd = "hello"
match cmd:
    case "quit":
        print("quitting")
    case "hello":
        print("greeting")
    case "help":
        print("showing help")
    default:
        print("unknown command")

```

Output: greeting

Match works with numbers, strings, booleans, and nil:

```

proc classify(n):
    match n:
        case 1:
            return "small"
        case 2:
            return "medium"
        default:
            return "big"

print(classify(1))    # small
print(classify(100))  # big

```

Cases are checked in order using value equality (**==**). Match is supported in the interpreter, C compiler (emits if/else-if chain), LLVM backend (emits branch instructions), and native codegen.

Part 5: Implementation Details and Runtime Behavior

5.1 Memory Layout and Object Representation

Stack vs. Heap: - **Stack-allocated** (cheap, temporary): `Value` struct for numbers, bools, nil, function pointers. - **Heap-allocated** (managed by GC): Arrays, dicts, tuples, strings, classes, instances, exceptions, generators.

Value Encoding (union-based):

```
Value v;
v.type = VAL_NUMBER;
v.as.number = 42.5;           // Efficient packing; only one field used per type

// Or:
v.type = VAL_ARRAY;
v.as.array = malloc(...);    // Pointer to ArrayValue struct
```

Implications: - Numbers and bools are always “cheap” to copy (no allocation). - Strings are immutable; reassignment creates new allocation (handled by parser/interpreter). - Arrays/dicts are mutable and passed by reference (modifications visible across scopes).

5.2 Call Stack and Environments

On Function Call: 1. **Create child environment:** `new_env = env_create(current_env)`. 2. **Bind parameters:** `env_define(new_env, param_name, param_value)`. 3. **Execute body:** `interpret(body, new_env)`. 4. **Return:** Pop environment, return to caller.

Nested Example:

```
Global: { x: 10 }
  ^- call add(3, 4)
      ^- Local: { a: 3, b: 4 } (parent: Global)
          ^- return a + b (7)
```

5.3 Exception Propagation

Raise Mechanism: - `raise "message"` or `raise exception_obj` → Sets `is_throwing = 1` in `ExecResult`. - `raise` converts non-string values: numbers become their string representation, booleans become “true”/“false”, nil becomes “nil”. - All statements/expressions check `is_throwing` and propagate up. - First matching `catch` clause executes; if none, exception exits function. - Sage uses generic catch (no type-based matching). All catches handle all exceptions.

Finally Guarantee: - `finally` block always runs, even if exception, return, break, or continue occurred in try/catch. - If `finally` raises an exception or returns, that overrides the try/catch result (matches Python/Java semantics). - If `finally` completes normally, the original try/catch result is preserved.

Memory Safety: - Exception messages are allocated outside the GC heap but tracked via `gc_track_external_allocation()`. - GC properly frees exception messages with `gc_track_external_free()` during collection. - `VAL_EXCEPTION` objects are marked during GC mark phase to prevent premature collection.

5.4 Generator State Management

Creation: - `proc gen_func(): yield 1; yield 2` parses as normal proc. - **Call** returns `VAL_GENERATOR` with: - **body:** Pointer to function body (`Stmt*`). - **closure:** Captured environment at call time. - **is_started:** 0 initially. - **is_exhausted:** 0 initially. - **current_stmt:** NULL (set on first `next()`).

Resumption: - **First next():** Initialize `gen_env` as child of closure, run body until first `yield`. - **Subsequent next():** Resume from `current_stmt` (set by prior `yield`), run until next `yield` or end. - **Yield execution:** Save `current_stmt = stmt->next`, return yielded value with `is_yielding = 1`.

Iteration Limits: - Direct `for-in` iteration over generators is explicitly rejected with a runtime error (Runtime Error: 'Generator' object is not iterable.). Always use `next()` for extraction.

Example Trace:

```
gen = my_gen()           # Create generator, is_started=0
next(gen)                # Initialize gen_env, run to first yield
                           # Yield 10; current_stmt = stmt after yield
next(gen)                # Resume from current_stmt, run to next yield
                           # Yield 20; current_stmt = stmt after yield
next(gen)                # Resume, reach end, is_exhausted=1
                           # Return nil
```

5.5 Operator Precedence and Associativity

Precedence (lowest to highest): 1. Assignment (=) 2. Logical OR (`or`) 3. Logical AND (`and`) 4. Bitwise OR (`|`) 5. Bitwise XOR (`^`) 6. Bitwise AND (`&`) 7. Equality (`==`, `!=`) 8. Comparison (`<`, `>`, `<=`, `>=`) 9. Shift (`<<`, `>>`) 10. Addition/Subtraction (+, -) 11. Multiplication/Division (*, /, %) 12. Unary (-, `not`, ~) * *Note: String indexing via bracket syntax (e.g., `s[i]`) returns a single-character string. Strings remain immutable. `slice(str, start, end)` and `indexof(str, sub)` are also available.* 13. Postfix (indexing, slicing, property access) 14. Primary (literals, variables, parens, function calls)

Associativity: - **Left:** Most binary operators (+, -, *, /, ==, !=, <, >, etc.). - **Right:** Assignment (=), unary (-).

5.6 Type Coercion and Truthiness

Truthiness (for `if/while` conditions): - **Falsy:** `nil`, `false`, 0, and empty strings (`""`). - **Truthy:** Everything else (empty arrays are `truthy`, objects, etc.).

Type Coercion: - **Addition:** If both are numbers, add; if either is string, concatenate; else `nil`. - **Comparison:** Numbers compared numerically; strings compared lexically; different types are unequal. - **String conversion:** `str()` native function converts any value to string.

Part 6: Compilation and Execution

6.1 Building the Interpreter

The repository exposes four build paths from source:

- Desktop build from C sources, which produces `sage` and `sage-lsp`
- Self-hosted bootstrap mode, which still builds `sage` first and then uses it to execute `src/sage/sage.sage`
- Pico/RP2040 builds through CMake and the Pico SDK
- SageMake: unified build system with auto-detection of platform, GPU (cuBLAS), NPU (NNAPI/SNPE/ONE), SIMD (NEON/RVV), and compiler backends

Desktop Build (Make):

```
make clean && make -j$(nproc)
./sage examples/hello.sage
```

ML Trainer Build:

```
make train-c          # Builds standalone C training binary (no Sage runtime); uses src/c/ml_backend.c dire
```

Desktop Build (CMake):

```
cmake -B build
cmake --build build
```

SageMake Build:

```
./sagemake info           # Show detected environment (GPU, NPU, SIMD, compiler)
./sagemake build          # Build sage interpreter
./sagemake chatbot --llvm # Compile chatbot via LLVM
./sagemake chatbot --c    # Compile via C backend
./sagemake chatbot --native # Compile via native asm
./sagemake train 200000 0.001 # Build trainer + train
./sagemake all            # Build everything
./sagemake --minimal build # Core only (no optional deps)

./sagemake info auto-detects GPU (cuBLAS), NPU (NNAPI/SNPE/ONE), SIMD (NEON/RVV), and
selects optimal compilation flags.
```

Self-Hosted Build / Bootstrap:

```
make sage-boot FILE=examples/hello.sage
make test-selfhost
```

```
# Or via CMake:
cmake -B build_sage -DBUILD_SAGE=ON
cmake --build build_sage
cmake --build build_sage --target test_selfhost
```

Pico Build:

```
cmake -B build_pico -DBUILD_PICO=ON -DPICO_BOARD=pico
cmake --build build_pico
```

Desktop builds require `libcurl` and `OpenSSL` development headers/libraries (e.g., `libcurl4-openssl-dev`) in addition to a C compiler, `make`, and/or `cmake`.

To build without network dependencies (such as `libcurl`), use `make CFLAGS_EXTRA="-DSAGE_NO_NET" SAGE_NO_NET=1`.

3 Build Parameter Reference Make Variables:

Variable	Default	Effect
CC	gcc	C compiler used by <code>make</code>
CFLAGS	-std=c11 -Wall -Wextra -Wpedantic -O2 -D_POSIX_C_SOURCE=200809L	Base compile flags for desktop builds
LDFLAGS	-lm -lpthread -ldl -lcurl -lssl -lcrypto	Desktop link flags; reduced to <code>-lm</code> when <code>PICO_BUILD</code> is set
DEBUG	0	<code>DEBUG=1</code> adds <code>-g -O0</code> <code>-DDEBUG</code>
PREFIX	/usr/local	Install prefix used by <code>make</code> <code>install</code>
FILE	unset	Required by <code>make</code> <code>sage-boot FILE=<path></code>
PICO_BUILD	unset	Internal Make switch that changes link flags for non-desktop builds

CMake Cache Variables / Environment Inputs:

Parameter	Default	Effect
BUILD_PICO	OFF	Enables Pico/RP2040 output and imports <code>pico_sdk_import.cmake</code> before <code>project()</code>
BUILD_SAGE	OFF	Enables bootstrap/self-hosted targets such as <code>sage_boot</code> and <code>test_selfhost</code>
ENABLE_DEBUG	OFF	Adds <code>-g -O0 -DDEBUG</code>
ENABLE_TESTS	OFF	Builds optional C test executables
CMAKE_BUILD_TYPE	generator default	Standard CMake build-type selector
CMAKE_C_COMPILER	toolchain default	Chooses the C compiler reported in the configuration summary
CMAKE_INSTALL_PREFIX	CMake default	Install destination for <code>cmake --install</code>
PICO_SDK_PATH	unset	Required for Pico builds unless already exported in the environment
PICO_BOARD	pico	Pico board name for SDK builds
SAGE_FILE	unset	Input path consumed by the <code>sage_boot</code> custom target

6.2 Main Entry Point (main.c)

`src/c/main.c` initializes the garbage collector, registers raw `argv` for the `sys` module, initializes the module cache, and then dispatches one of the top-level modes below.

3 sage CLI Parameter Reference

Command	Meaning	Notes
<code>sage</code>	Start the interactive REPL	Same as <code>sage --repl</code>
<code>sage --repl</code>	Start the interactive REPL	Supports multi-line blocks and recovery after REPL errors
<code>sage --help</code>	Print usage text	Covers compiler, tooling, and REPL entry points
<code>sage -c "source"</code>	Execute a source string	No file is loaded
<code>sage <file.sage> [arg ...]</code>	Execute a Sage file	Additional CLI arguments are visible through <code>sys.args()</code>
<code>sage <file.sgvm></code>	Execute an SGVM binary	Runs via integrated MetalVM engine
<code>sage --lsp</code>	Start the LSP server on stdin/stdout	<code>sage-lsp</code> is the standalone wrapper binary
<code>sage fmt <file></code>	Format a file in place	Prints Formatted: <code><file></code> on success

Command	Meaning	Notes
<code>sage fmt --check <file></code>	Check formatting without rewriting	Exit code 1 when changes are needed
<code>sage lint <file></code>	Run the static linter	Exit code 1 when issues are found

Compiler Command	Default Output	Options
<code>sage --emit-vm <input.sage></code>	<code><input>.svm</code>	<code>-o <path>, -00, -01, -02, -03, -g</code>
<code>sage --sgvm <input.sage></code>	<code><input>.sgvm</code>	<code>-o <path>, -00, -01, -02, -03, -g</code>
<code>sage --compile <input.sage></code>	<code><input-without-.sage></code>	<code>-o <path>, --cc <compiler>, -00, -01, -02, -03, -g</code>
<code>sage --emit-llvm <input.sage></code>	<code><input>.ll</code>	<code>-o <path>, -00, -01, -02, -03, -g</code>
<code>sage --compile-llvm <input.sage></code>	<code><input-without-.sage></code>	<code>-o <path>, -00, -01, -02, -03, -g</code>
<code>sage --emit-asm <input.sage></code>	<code><input>.s</code>	<code>-o <path>, --target <arch[-profile]>, -00, -01, -02, -03, -g</code>
<code>sage --compile-native <input.sage></code>	hosted: <code><input-without-.sage></code> ; non-hosted profiles: <code><input-without-.sage>.o</code>	<code>-o <path>, --target <arch[-profile]>, -00, -01, -02, -03, -g</code>
<code>sage --emit-pico-c <input.sage></code>	<code><input>.pico.c</code>	<code>-o <path></code>
<code>sage --compile-pico <input.sage></code>	<code>.tmp/<program-name> plus <program-name>.uf2</code>	<code>-o <dir>, --board <name>, --name <program>, --sdk <path>, --chip <type></code>
<code>--chip <type></code>	<code>--compile-pico</code>	Chip type (rp2040, rp2350-arm, or rp2350-riscv); defaults to rp2040
<code>sage --compile-bare <input.sage></code>	<code><input-without-.sage>.elf</code>	<code>-o <path>, --target <arch>, -00, -01, -02, -03, -g</code>
<code>sage --compile-uefi <input.sage></code>	<code><input-without-.sage>.efi</code>	<code>-o <path>, --target x86_64\ aarch64, -00, -01, -02, -03, -g</code>

Compiler Command	Default Output	Options
<code>sage --jit <input.sage></code>	JIT executable (temporary)	<code>-o <path></code> for self-extracting executable with module bundling

Option	Applies To	Meaning
<code>-o <path></code>	All emit/compile commands, including <code>--compile-pico</code>	Output file or build directory
<code>--cc <compiler></code>	<code>--compile</code>	Overrides the host C compiler; defaults to <code>cc</code>
<code>--target <arch[-profile]></code>	<code>--emit-asm</code> , <code>--compile-native</code>	Base targets: <code>x86-64</code> , <code>x86_64</code> , <code>aarch64</code> , <code>arm64</code> , <code>rv64</code> , <code>riscv64</code> , <code>mips</code> , <code>mips32</code> , <code>mips74k</code> ; profile suffixes: <code>-baremetal</code> , <code>-osdev</code> , <code>-uefi</code>
<code>-O0 / -O1 / -O2 / -O3</code>	C, LLVM, and native codegen commands	Selects the optimization pass level
<code>-g</code>	C, LLVM, asm, and native codegen commands	Enables debug information in generated output
<code>--math-work=MODES</code>	All runtime modes	Comma-separated list of default math visualization formats: <code>grade</code> , <code>exec</code> , <code>bitwise</code>
<code>--board <name></code>	<code>--compile-pico</code>	Pico board name; defaults to <code>pico</code>
<code>--name <program></code>	<code>--compile-pico</code>	Overrides the generated program name derived from the input file
<code>--sdk <path></code>	<code>--compile-pico</code>	Pico SDK path; falls back to <code>PICO_SDK_PATH</code>

LLVM backend notes:

- `--compile-llvm` produces a fully native binary by emitting LLVM IR and invoking `clang` to compile and link it. The result is a standalone executable with no Sage runtime dependency.
- As a real-world data point, the SageLLM chatbot (`models/chatbots/sagellm_chatbot.sage`) compiles to a **124 KB standalone binary** via `--compile-llvm`, including the full inference loop and tokenizer.

Profile notes for native backend:

- `hosted` (no suffix): default behavior.
- `-baremetal / -osdev`: emits `sage_entry` and freestanding object-oriented output.
- `-uefi`: emits `efi_main`; current stage emits a freestanding object, with full PE/COFF image linking planned as follow-up work.

6.3 Program Execution Flow

1. **Read source** from file.
2. **Initialize lexer** with source string.
3. **Initialize parser** (advance to first token).

4. **Create global environment** and initialize `stdlib`.
5. **Parse declarations** (statements and function definitions).
6. **Execute each statement** using the selected runtime backend.
7. **Return** at EOF.

6.4 Runtime Backends

The C-hosted `sage` binary now supports several runtime selections:

Mode	Command	Current Behavior
<code>ast</code>	<code>sage --runtime ast</code> <code>file.sage</code>	Original tree-walking interpreter; highest maturity and easiest to debug
<code>bytecode</code>	<code>sage --runtime bytecode</code> <code>file.sage</code>	Lowers each top-level statement to bytecode and executes it on the stack VM
<code>auto</code>	<code>sage --runtime auto</code> <code>file.sage</code>	Default; resolves to JIT on hosted, AST on bare-metal
<code>jit</code>	<code>sage --runtime jit</code> <code>file.sage</code>	Enables JIT profiling and type feedback
<code>aot</code>	<code>sage --runtime aot</code> <code>file.sage</code>	Enables Ahead-of-Time type specialization

Current VM architecture:

- The lexer and parser are unchanged; the VM reuses the same AST front-end.
- Each parsed top-level statement is compiled to a transient bytecode chunk, then executed immediately.
- The C-hosted toolchain can also emit a strict ahead-of-time VM artifact with `sage --emit-vm file.sage` and execute it later with `sage --run-vm file.svm`.
- The self-hosted CLI can emit the same artifact format with `sage sage.sage --emit-vm file.sage`, including compiled proc bodies and returns.
- `sage --runtime bytecode` remains **hybrid** today: unsupported top-level statements fall back to the AST interpreter through an explicit AST-bridge opcode instead of failing the whole program.
- `sage --emit-vm` is intentionally stricter: unsupported constructs fail compilation instead of silently bridging at runtime.
- Values, environments, modules, classes, instances, and the GC are shared between AST and bytecode execution.

What runs natively in the VM today:

- Literals (`number`, `string`, `bool`, `nil`)
- Global variables (`let`, assignment, reads)
- Arithmetic, comparison, logical, and bitwise operators
- Arrays, tuples, dicts, indexing, slicing, and property access
- `print`, expression statements, `if`, `while`, and array `for`
- `break` and `continue` inside `while/for` loops (compiled natively with loop context stack)
- Ahead-of-time proc definitions, proc calls, nested proc calls, explicit `return`, and implicit `nil` returns
- Calls to native functions, Sage functions, classes, and instance methods
- 30 GPU hot-path opcodes for frame-loop performance (`poll_events`, `key_pressed`, `cmd_draw`, etc.)

What still bridges or stays unsupported:

- In hybrid `--runtime bytecode` mode: class definitions, module imports, exception handling (`try/catch/raise`), `defer`, `match`, and `yield` fall back to the AST interpreter via `BC_OP_EXEC_AST_STMT`. Opcodes are defined for future native support (`BC_OP_CLASS`, `BC_OP_IMPORT`, `BC_OP_SETUP_TRY`, `BC_OP_RAISE`, etc.).
- In strict `--emit-vm` mode: these constructs fail compilation instead of bridging.

Security: The VM validates all constant pool accesses (`VM_CHECK_CONST`) and AST statement indices (`VM_CHECK_AST`) to prevent buffer overflow from malformed bytecode. Stack depth is bounded at 1024 entries. All memory allocation uses OOM-safe wrappers.

The practical result is that `bytecode` mode is already useful for long-running scripts and engine-style workloads, `--emit-vm` is now a real ahead-of-time path for a meaningful strict subset of Sage, and `ast` mode remains the reference path for maximum behavioral confidence.

Part 7: Advanced Topics and Internals

7.1 Memory Leaks and GC Limitations

Known Limitations: - **String interning:** Strings are not interned; identical strings create separate allocations. - **Non-GC allocations:** AST nodes, token buffers, and other manual allocations outside the `Value` heap still need explicit ownership and cleanup discipline. - **Manual `malloc`:** Some internal allocations (e.g., AST nodes, Token data) use `malloc` without GC tracking. Only heap-allocated `Values` are GC'd.

Cycle collection:

- The GC is now a tracing mark-and-sweep collector that can reclaim **circular references** among GC-managed Sage values such as arrays, dicts, tuples, classes, instances, functions, and generators.
- Cycles are only collectable when every object in the cycle is represented in the managed `Value` graph. Raw pointers, external allocations, and ad hoc native-side ownership still sit outside that guarantee.

LLVM Backend Loop Variable Limitation:

- **For-loop iteration variables cannot be modified to simulate a break in LLVM-compiled binaries.** Code like `j = len(arr)` inside a for loop has no effect on the loop control variable in `--compile-llvm` output. Always use `break` to exit loops early. This works correctly in the interpreter and all other compiler backends (C, native).

Practical Tips: - Call `gc_collect()` periodically in long-running loops to prevent heap explosion. - Avoid creating large temporary structures in loops. - Reuse arrays/dicts by modifying in-place rather than creating new ones.

7.2 Performance Characteristics

Interpreter Overhead: - **AST runtime:** Walks the parsed tree directly. It is still the simplest and most mature execution path. - **Bytecode VM:** Adds a statement-at-a-time lowering step, then executes stack bytecode. On mixed arithmetic/loop workloads it is already modestly faster than AST mode. - **Threaded Dispatch:** The VM uses computed gotos (threaded interpretation) for high-performance instruction dispatch on supported compilers, significantly reducing loop overhead. - **Register Spilling:** The VM stack pointer and instruction pointer are kept in CPU registers where possible, minimizing memory traffic during the execution loop. - **Nested scopes:** Linear search up parent chain; $O(\text{depth})$ for variable lookup. - **GC pauses:** Mark-and-sweep still pauses execution during collection, but the trigger is now dynamic and based on both object count and managed bytes.

Optimization Opportunities: - Bytecode functions and method bodies, so fewer calls bridge back to AST. - Local-slot or register-based VM frames instead of env-backed global lookups for more code. - Variable lookup caching (memoization). - Generational GC (mark only young objects frequently). - JIT compilation for hot paths.

3 Current Recipe Benchmark

The repository now includes a five-recipe benchmark:

```
python3 scripts/benchmark_recipes.py --runs 5 --warmups 1
```

It measures:

- `sage-interpreted-c-ast`: the C-hosted tree-walking interpreter
- `sage-interpreted-vm`: the C-hosted bytecode VM runtime
- `sage-compiled-c`: the native C backend compiled to a host executable
- `sage-compiled-vm`: the C-hosted ahead-of-time VM artifact (`--emit-vm + --run-vm`)
- `sage-compiled-sage`: the self-hosted Sage compiler path that emits C and then compiles it with the host toolchain

Workload source: `benchmarks/runtime_compare.sage`

Generated chart assets:

- `assets/charts/benchmark-recipes-total.svg`
- `assets/charts/benchmark-recipes-run.svg`

Current caveats:

- `sage-compiled-vm` now passes the default benchmark workload and is charted as a real execution lane, but it is still a strict subset backend: unsupported constructs fail during `--emit-vm` instead of falling back at runtime.
- `sage-compiled-sage` is still experimental and currently fails checksum validation on the default benchmark workload, so it is called out in the chart footer instead of being charted as a valid timing result.

Interpretation:

- `sage-interpreted-vm` is the direct runtime-backend comparison against `sage-interpreted-c-ast`.
- `sage-compiled-c` has the lowest execution-only runtime on the default workload, but its total wall time includes code generation and host compilation.
- The total-time chart answers “time to result”; the execution-only chart answers “steady-state runtime after the binary already exists.”

3 Sage vs Python 3 Benchmarks

A separate benchmark suite compares all Sage execution paths against CPython 3.x:

```
make benchmark-python      # Console table output
make benchmark-python-md   # Markdown table output
```

10 paired workloads (`benchmarks/01_fibonacci.sage + .py` through `benchmarks/10_primes_sieve`):

Benchmark	What it tests
<code>fibonacci</code>	Recursive function call overhead
<code>loop_sum</code>	Raw loop + arithmetic throughput
<code>string_concat</code>	String allocation + GC pressure
<code>array_ops</code>	Dynamic array growth + iteration
<code>dict_ops</code>	Hash table insert + lookup
<code>class_method</code>	OOP method dispatch
<code>nested_loops</code>	Control flow with break/continue
<code>exception_handling</code>	Try/catch overhead
<code>recursion_closures</code>	Closure creation + higher-order calls
<code>primes_sieve</code>	Array access + conditional logic

Five recipes tested: Python 3, Sage AST, Sage VM, Sage C (compiled), Sage LLVM (compiled).

Typical performance characteristics: - **LLVM compiled**: 2-8x faster than CPython (fibonacci 3.9x, loop sum 7.7x) - **C compiled**: 2-4x faster than CPython on most workloads - **AST interpreter**: On par with CPython (faster on string/array ops, slower on deep recursion) - **Bytecode VM**: Similar to AST, with measurable speedups on tight loops

7.3 Extending SageLang with Native Functions

Adding a Native Function:

```
// In interpreter.c, add:
static Value my_function(int argCount, Value* args) {
    if (argCount != 2) return val_nil();
    if (!IS_NUMBER(args[0]) || !IS_NUMBER(args[1])) return val_nil();
    double result = AS_NUMBER(args[0]) + AS_NUMBER(args[1]) * 2;
    return val_number(result);
}

// In init_stdlib(), add:
env_define(env, "my_func", 7, val_native(my_function));
```

Usage from SageLang:

```
print my_func(3, 4) # 3 + 4*2 = 11
```

7.4 Debugging and Introspection

Print Value:

```
void print_value(Value v); // In value.c
```

GC Stats:

```
let stats = gc_stats()
print stats # {"bytes_allocated": ..., "num_objects": ..., ...}
```

Tracing Execution (manual): - Add `printf()` statements to `interpreter.c` before/after statement/expression evaluation. - Mark garbage objects: Modify `gc_sweep()` to not free unmarked objects, inspect them.

7.5 Interpreter Safety Limits

SageLang enforces several compile-time and runtime limits to prevent crashes from malicious or malformed input:

Limit	Constant	Value	Location	Behavior on Violation
Recursion depth (logical)	<code>MAX_RECURSION_DEPTH</code>	2000	<code>interpreter.c</code>	Catchable exception
Recursion depth (real stack)	<code>stack_danger()</code>	75% of <code>RLIMIT_STACK</code>	<code>interpreter.c</code>	Catchable exception — fires before the OS stack is exhausted regardless of frame size
Parser nesting	<code>MAX_PARSER_DEPTH</code>	100000	<code>parser.c</code>	Parse error
Loop iterations	<code>MAX_LOOP_ITERATIONS</code>	500,000	<code>interpreter.c</code>	Catchable exception
String literal length	<code>MAX_STRING_LENGTH</code>	4096	<code>lexer.c</code>	Parse error
Function arguments	Stack array	255	<code>interpreter.c</code>	Runtime error

The stack-proximity guard (v4.1.16) tracks a per-thread stack origin and refuses recursion once within the safety margin of the real OS limit, so deep or fat-framed call chains fail with a catchable exception instead

of SIGSEGV. At startup the runtime also raises the soft stack limit to 512 MB when the hard limit permits, giving self-hosted double-interpretation chains headroom. Since v4.2.9 the margin is computed from the **calling thread's actual stack bounds** rather than from `RLIMIT_STACK`: the raised limit only describes the main thread, while worker threads are created with a fixed stack that a later `setrlimit` cannot change. Deriving the budget from the rlimit therefore overstated a worker's headroom (~384 MB for an 8 MB stack) and let runaway recursion in any thread segfault. Each thread now gets a budget for the stack it really has; on a typical Linux host a worker tops out cleanly at roughly 200 levels of Sage recursion, while the main thread still reaches the 12000-level counter limit.

Null function guards: The interpreter checks for null pointers in both `VAL_FUNCTION` and `VAL_NATIVE` call paths before dispatch. A null callee produces a runtime error and returns `nil`.

Type-safe accessor macros (`value.h`):

Macro	Returns on type mismatch
<code>SAGE_AS_STRING(v)</code>	<code>""</code> (empty string)
<code>SAGE_AS_NUMBER(v)</code>	<code>0.0</code>
<code>SAGE_AS_BOOL(v)</code>	<code>0</code> (false)

These are intended for native function implementations that want defensive coercion. The standard `IS_*/AS_*` pattern (check type, then access) remains the recommended approach for most code:

```
static Value my_native(int argCount, Value* args) {
    if (argCount < 1 || !IS_STRING(args[0])) return val_nil();
    const char* s = AS_STRING(args[0]); // safe - guarded by IS_STRING above
    // ...
}
```

Part 8: Future Directions and Design Notes

8.1 Planned Features (Incomplete Implementations)

Pattern Matching (`match/case`): Fully implemented across all backends (see Section 4.9).

Defer Statements: Fully implemented with LIFO ordering and scope-exit semantics (see Section 4.8).

Multiple Inheritance / Traits: - Currently single-parent inheritance only.

Operator Overloading: - No support for custom `__add__()`, etc.

Type Annotations: - Fully supported for variables (`let x: Int = 42`) and procedures (`proc add(a: Int, b: Int) -> Int:`). - Dotted names are supported for qualified types (e.g., `let x: vfs.VFS`). - Type checking is enforced during a static analysis pass (`pass_typecheck` in `typecheck.c`) prior to compilation or interpretation.

8.2 Design Decisions and Rationale

Why Indentation-Based Syntax? - Python familiarity for embedded systems developers. - Reduces bracket clutter in embedded/IoT code. - Lexer complexity worth the syntactic clarity.

Why Mark-and-Sweep Over Reference Counting? - Simpler to implement; no cycle detection overhead. - Acceptable for embedded contexts where pause is tolerable. - Future: Generational GC for better pause times.

Why Tree-Walking Interpreter? - Rapid prototyping and debugging. - Minimal complexity; easy to extend. - Suitable for small programs; not for performance-critical workloads.

Why Classes Over Functions-as-Objects? - Explicit OOP model matches C++ style (familiar to embedded developers). - Clearer semantics for **self** and method dispatch.

Part 9: Bundled Library Modules (lib/)

SageLang ships with a set of **pure-Sage library modules** in the `lib/` directory. These are importable via `import` or `from ... import` when running from the project root (the module resolver searches `./lib/`).

Note: `lib/math.sage` and native `math` share a name — the native module takes precedence. Use `import math` for native math functions (`sin`, `cos`, `sqrt`, etc.) and `from math import ...` for native `math`. The `lib/math.sage` helpers (`factorial`, `gcd`, etc.) are available if no native module shadows them.

9.1 Arrays Module (`lib/arrays.sage`)

Functional array utilities:

```
from arrays import map, filter, reduce, reverse, unique, flatten, zip, chunk
```

```
proc double(x):
    return x * 2
```

```
print map([1, 2, 3], double)      # [2, 4, 6]
print filter([1, 2, 3, 4], is_even) # [2, 4]
print reverse([1, 2, 3])          # [3, 2, 1]
print unique([1, 2, 2, 3])         # [1, 2, 3]
print flatten([[1, 2], [3, 4]])    # [1, 2, 3, 4]
print zip([1, 2], ["a", "b"])      # [(1, a), (2, b)]
print chunk([1, 2, 3, 4, 5], 2)    # [[1, 2], [3, 4], [5]]
```

Function	Description
<code>copy(arr)</code>	Shallow copy
<code>append_all(target, extra)</code>	Extend target in-place
<code>concat(a, b)</code>	Return new merged array
<code>reverse(arr)</code>	Return reversed copy
<code>map(arr, fn)</code>	Apply fn to each element
<code>filter(arr, pred)</code>	Keep elements where pred is true
<code>reduce(arr, init, fn)</code>	Fold left with accumulator
<code>contains(arr, val)</code>	Check membership
<code>index_of(arr, val)</code>	Find index (-1 if missing)
<code>find(arr, pred)</code>	First element matching predicate
<code>unique(arr)</code>	Remove duplicates
<code>flatten(nested)</code>	Flatten one level of nesting
<code>take(arr, n)</code>	First n elements
<code>drop(arr, n)</code>	Skip first n elements
<code>zip(a, b)</code>	Pair elements as tuples
<code>chunk(arr, size)</code>	Split into chunks

9.2 Strings Module (`lib/strings.sage`)

String manipulation utilities:

```
from strings import words, compact, pad_left, dash_case, from_bin
```

```

print words(" hello world ") # [hello, world]
print compact(" hello world ") # hello world
print pad_left("42", 5, "0") # 00042
print dash_case("hello world") # hello-world
print from_bin("0b1010") # 10

```

Function	Description
<code>words(text)</code>	Split on whitespace, remove empties
<code>compact(text)</code>	Collapse whitespace to single spaces
<code>contains(text, sub)</code>	Substring check
<code>count_substring(text, sub)</code>	Count occurrences
<code>repeat(text, n)</code>	Repeat string n times
<code>pad_left(text, width, pad)</code>	Left-pad to width
<code>pad_right(text, width, pad)</code>	Right-pad to width
<code>surround(text, left, right)</code>	Wrap with delimiters
<code>csv(values)</code>	Join with commas
<code>dash_case(text)</code>	Convert to dash-case
<code>snake_case(text)</code>	Convert to snake_case
<code>endswith(text, suffix)</code>	Check suffix
<code>from_bin(bits)</code>	Binary string to number

9.3 Dicts Module (`lib/dicts.sage`)

Dictionary query and manipulation helpers:

```
from dicts import has, size, get_or, has_all, entries, remove_keys
```

```

let d = {}
d["x"] = 10
d["y"] = 20

```

```

print has(d, "x") # true
print size(d) # 2
print get_or(d, "z", "default") # default
print has_all(d, ["x", "y"]) # true

```

Function	Description
<code>keys(d)</code>	Wrapper for <code>dict_keys</code>
<code>values(d)</code>	Wrapper for <code>dict_values</code>
<code>size(d)</code>	Number of keys
<code>has(d, key)</code>	Key existence check
<code>get_or(d, key, fallback)</code>	Get with default
<code>entries(d)</code>	Array of (key, value) tuples
<code>has_all(d, keys)</code>	All keys present?
<code>has_any(d, keys)</code>	Any key present?
<code>select_values(d, keys, fallback)</code>	Get multiple values
<code>remove_keys(d, keys)</code>	Delete keys in-place
<code>count_missing(d, keys)</code>	Count absent keys

9.4 Iter Module (`lib/iter.sage`)

Reusable generator functions:

```

from iter import count, range_step, cycle, enumerate_array, take

let evens = range_step(0, 20, 2)
print take(evens, 5) # [0, 2, 4, 6, 8]

let c = cycle([1, 2, 3])
print take(c, 7) # [1, 2, 3, 1, 2, 3, 1]

let e = enumerate_array(["a", "b", "c"])
print next(e) # (0, a)
print next(e) # (1, b)

```

Function	Description
<code>count(start, step)</code>	Infinite counter generator
<code>range_step(start, end, step)</code>	Range with custom step
<code>repeat(value, count)</code>	Yield value n times
<code>repeat_forever(value)</code>	Infinite repeat
<code>enumerate_array(arr)</code>	Yield (index, value) tuples
<code>cycle(arr)</code>	Infinite cycle through array
<code>take(gen, n)</code>	Collect n values from generator
<code>nth(gen, index)</code>	Get nth value from generator

9.5 Stats Module (lib/stats.sage)

Statistical functions:

```

from stats import mean, variance, stddev, cumulative, normalize

print mean([1, 2, 3, 4, 5]) # 3
print variance([1, 2, 3, 4, 5]) # 2
print cumulative([1, 2, 3, 4, 5]) # [1, 3, 6, 10, 15]
print normalize([1, 2, 3, 4, 5]) # [0, 0.25, 0.5, 0.75, 1]

```

Function	Description
<code>sum(values)</code>	Sum of array
<code>product(values)</code>	Product of array
<code>min_value(values)</code>	Minimum element
<code>max_value(values)</code>	Maximum element
<code>mean(values)</code>	Arithmetic mean
<code>range_span(values)</code>	max - min
<code>cumulative(values)</code>	Running totals
<code>variance(values)</code>	Population variance
<code>stddev(values)</code>	Standard deviation
<code>normalize(values)</code>	Scale to [0, 1]

9.6 Utils Module (lib/utils.sage)

General-purpose helpers:

```

from utils import default_if_nil, choose, between, head, last, repeat_value

print default_if_nil(nil, 42) # 42
print choose(true, "yes", "no") # yes

```

```

print between(5, 1, 10)      # true
print head([1, 2, 3])       # 1
print last([1, 2, 3])       # 3
print repeat_value(0, 5)    # [0, 0, 0, 0, 0]

```

Function	Description
<code>identity(x)</code>	Return x unchanged
<code>choose(cond, a, b)</code>	Ternary selection
<code>default_if_nil(val, fallback)</code>	Nil coalescing
<code>is_even(n) / is_odd(n)</code>	Parity check
<code>between(val, lo, hi)</code>	Range check
<code>swap(a, b)</code>	Return (b, a) tuple
<code>head(arr) / last(arr)</code>	First/last element (nil if empty)
<code>repeat_value(val, n)</code>	Array of n copies
<code>times(n, fn)</code>	Call fn(i) n times

9.7 Assert Module (`lib/assert.sage`)

Test assertion helpers (raise on failure):

```
from assert import assert_equal, assert_true, assert_close
```

```

assert_true(len([1, 2, 3]) == 3, "wrong length")
assert_equal(2 + 2, 4, "math is broken")
assert_close(3.14, 3.14159, 0.01, "not close enough")

```

Function	Description
<code>fail(msg)</code>	Always raise
<code>assert_true(cond, msg)</code>	Raise if false
<code>assert_false(cond, msg)</code>	Raise if true
<code>assert_equal(actual, expected, msg)</code>	Raise if not equal
<code>assert_nil(val, msg)</code>	Raise if not nil
<code>assert_not_nil(val, msg)</code>	Raise if nil
<code>assert_close(actual, expected, tol, msg)</code>	Raise if difference > tolerance
<code>assert_array_contains(arr, val, msg)</code>	Raise if val not in arr

9.8 Math Module (`lib/math.sage`)

Pure-Sage math helpers (shadowed by native `math` module — use when native module is unavailable):

Function	Description
<code>add, sub, mul, div</code>	Basic arithmetic
<code>min, max, abs, sign</code>	Comparison helpers
<code>clamp(val, lo, hi)</code>	Clamp to range
<code>square(x), cube(x)</code>	Powers
<code>lerp(a, b, t)</code>	Linear interpolation
<code>pow_int(base, exp)</code>	Integer exponentiation
<code>factorial(n)</code>	n!
<code>gcd(a, b), lcm(a, b)</code>	GCD and LCM
<code>sum(arr), product(arr), mean(arr)</code>	Aggregates

Function	Description
<code>sqrt(n)</code>	Newton's method approximation
<code>printm(expr, backend, formats)</code>	Evaluate expression and show work
<code>distance(x1, y1, x2, y2)</code>	Euclidean distance

Arithmetic Visualization (`printm`):

The `math.printm()` function evaluates an arithmetic expression provided as a string and visualizes the step-by-step “work” for each operation.

```
import math
math.printm("123 + 456 * 2", backend="sage")
```

- **backend:**
 - “sage” (default): Pure SageLang implementation using vertical “grade-school” format.
 - “c”: Uses native C-level arithmetic bridges.
 - “asm”: Uses inline assembly (x86-64, aarch64, rv64, mips) and calls `printf` via FFI to show the instruction-level execution.
- **formats:** An array of strings to customize output.
 - “grade”: Vertical grade-school arithmetic visualization.
 - “exec”: Show the underlying execution mode (e.g., “[ASM] Using ADD instruction”).
 - “bitwise”: Reserved for bitwise operation visualization.

The default format can be configured globally using the `--math-work` CLI flag. | `normalize(val, lo, hi)`
| Scale to [0, 1] |

9.9 GPU Libraries

SageLang ships with 18 GPU/rendering library modules in `lib/graphics/`. All are imported with the `graphics.` prefix (e.g., `import graphics.vulkan` binds as `vulkan`):

Module	Import	Lines	Purpose
<code>vulkan.sage</code>	<code>import graphics.vulkan</code>	~425	Ergonomic Vulkan builder API (string-based buffer/shader/pipeline creation)
<code>gpu.sage</code>	<code>import graphics.gpu</code>	~150	High-level compute helpers (one-shot dispatch, ping-pong buffers)
<code>opengl.sage</code>	<code>import graphics.opengl</code>	~400	Drop-in OpenGL backend (same API as <code>gpu</code> module, OpenGL 4.5 init)
<code>ui.sage</code>	<code>import graphics.ui</code>	~400	Immediate-mode GUI widgets (windows, buttons, sliders, menus, text inputs)
<code>math3d.sage</code>	<code>import graphics.math3d</code>	~250	<code>vec2/3/4</code> , <code>mat4</code> , perspective/ortho projections, camera utilities
<code>mesh.sage</code>	<code>import graphics.mesh</code>	~300	Procedural geometry (cube, plane, sphere), OBJ loading, GPU upload
<code>renderer.sage</code>	<code>import graphics.renderer</code>	~200	Frame loop management (depth buffer, render pass, sync)

Module	Import	Lines	Purpose
<code>material.sage</code>	<code>import graphics.material</code>	~150	Shader+texture+descriptor binding, material presets
<code>scene.sage</code>	<code>import graphics.scene</code>	~200	Node hierarchy, transforms, traversal, <code>find_by_name</code>
<code>pbr.sage</code>	<code>import graphics.pbr</code>	~300	Cook-Torrance PBR materials, point/directional lights, IBL
<code>postprocess.sage</code>	<code>import graphics.postprocess</code>	~250	HDR targets, bloom chain, tone mapping (ACES/Reinhard)
<code>shadows.sage</code>	<code>import graphics.shadows</code>	~200	Shadow maps, depth-only passes, cascade shadow maps
<code>deferred.sage</code>	<code>import graphics.deferred</code>	~300	G-buffer (4 MRT), SSAO (32 samples), SSR (64-step raymarch)
<code>taa.sage</code>	<code>import graphics.taa</code>	~150	Temporal anti-aliasing (Halton jitter, history blend)
<code>gltf.sage</code>	<code>import graphics.gltf</code>	~200	glTF 2.0 JSON loading, mesh/material extraction
<code>asset_cache.sage</code>	<code>import graphics.asset_cache</code>	~100	Shader/texture/mesh caching and deduplication
<code>frame_graph.sage</code>	<code>import graphics.frame_graph</code>	~150	Pass dependency ordering (topological sort)
<code>debug_ui.sage</code>	<code>import graphics.debug_ui</code>	~150	FPS tracking, custom debug values, toggle overlay

9.10 OS Development Libraries

SageLang ships with 52 OS/bare-metal development modules across `lib/os/`, `lib/os/boot/`, `lib/os/kernel/`, and `lib/os/image/`:

Module	Import	Purpose
<code>fat.sage</code>	<code>import os.fat</code>	FAT8/12/16/32 boot sector parser, cluster-to-LBA, FAT entry offsets
<code>fat_dir.sage</code>	<code>import os.fat_dir</code>	FAT directory traversal, file reading, path resolution, cluster chains
<code>elf.sage</code>	<code>import os.elf</code>	ELF32/64 header, program/section headers, string table lookup
<code>mbr.sage</code>	<code>import os.mbr</code>	MBR partition table, CHS decode, bootable partition finder
<code>gpt.sage</code>	<code>import os.gpt</code>	GPT header, GUID parsing, partition type identification
<code>pe.sage</code>	<code>import os.pe</code>	PE/COFF binary parser, DOS/COFF/optional headers, UEFI app detection
<code>pci.sage</code>	<code>import os.pci</code>	PCI config space (Type 0/1), BAR decode, capability lists
<code>uefi.sage</code>	<code>import os.uefi</code>	EFI memory map, config tables, RSDP, ACPI SDT headers

Module	Import	Purpose
<code>tmpfs.sage</code>	<code>import os.tmpfs</code>	In-memory filesystem implementation
<code>sync.sage</code>	<code>import os.sync</code>	Synchronization primitives for the OS
<code>cgroups.sage</code>	<code>import os.linux.cgroups</code>	Linux Control Groups API for resource limits
<code>namespace.sage</code>	<code>import os.linux.namespace</code>	Linux Namespaces API for containerization
<code>netlink.sage</code>	<code>import os.linux.netlink</code>	Linux Netlink socket bindings
<code>core.sage</code>	<code>import metal.core</code>	Bare-metal CPU and architecture-specific routines (fences, yields, spinlocks)
<code>timer.sage</code>	<code>import metal.timer</code>	Hardware timer control (periodic/oneshot tracking)
<code>acpi.sage</code>	<code>import os.acpi</code>	MADT (APIC), FADT, HPET, MCFG parsers, processor enumeration
<code>paging.sage</code>	<code>import os.paging</code>	x86-64 page tables, PTE flags, identity/higher-half mapping helpers
<code>idt.sage</code>	<code>import os.idt</code>	x86-64 IDT gate construction, exception vectors, PIC remapping
<code>serial.sage</code>	<code>import os.serial</code>	UART/COM port configuration, init sequences, debug output encoding
<code>dtb.sage</code>	<code>import os.dtb</code>	Flattened Device Tree parser for ARM64/RISC-V (nodes, properties, search)
<code>alloc.sage</code>	<code>import os.alloc</code>	Bump, free-list, and bitmap page allocators for kernel memory
<code>vfs.sage</code>	<code>import os.vfs</code>	Virtual filesystem layer with mount table, path utilities, memfs backend

Module	Import	Purpose
<pre> ext.sage" import os.ext ext2/3/4 superblock, inode table, directory entries, extent tree btrfs.sage import os.btrfs Btrfs superblock, chunk tree, root tree, subvolumes, checksums f2fs.sage import os.f2fs F2FS superblock, checkpoint, segment info, node/data addressing boot/multiboot.sage import os.boot.multiboot Multiboot2 header generation, tag building, boot info parsing boot/gdt.sage import os.boot.gdt x86_64 GDT descriptor construction, TSS entries, LGDT sequence boot/start.sage import os.boot.start x86_64 startup assembly generation (long mode entry, stack setup) boot/linker.sage import os.boot.linker Linker script generation for bare-metal ELF kernels boot/bios.sage import os.boot.bios BIOS interrupt interface (INT 0x10, 0x13, 0x15, 0x16) for legacy x86 boot boot/e820.sage import os.boot.e820 BIOS memory map collection and normalization boot/a20.sage import os.boot.a20 A20 gate enabling methods (BIOS, Fast A20, Keyboard Controller) boot/cpuid.sage import os.boot.cpuid CPU feature detection (Long mode, MSR, SSE, Vendor/Brand strings) boot/real_mode.sage import os.boot.real_mode Helpers for segment:offset addressing and real-mode stack setup boot/prot_mode.sage import os.boot.prot_mode Transition logic from real mode to 32-bit protected mode boot/long_mode.sage import os.boot.long_mode Transition logic from </pre>	Linux /sys filesystem interface	

9.11 Networking Libraries

SageLang ships with 8 high-level networking modules in `lib/net/`, built on top of the native `socket`, `tcp`, `http`, and `ssl` modules:

Module	Import	Purpose
<code>url.sage</code>	<code>import net.url</code>	URL parsing/building, percent-encoding/decoding, query strings
<code>headers.sage</code>	<code>import net.headers</code>	HTTP header parsing/building, content-type inspection, constants
<code>request.sage</code>	<code>import net.request</code>	HTTP request builder with fluent API, auth helpers, status utilities
<code>server.sage</code>	<code>import net.server</code>	TCP/HTTP server framework with routing, request parsing, response builders
<code>websocket.sage</code>	<code>import net.websocket</code>	WebSocket frame building/parsing (RFC 6455), upgrade handshake
<code>mime.sage</code>	<code>import net.mime</code>	MIME type lookup from file extensions (80+ types)
<code>dns.sage</code>	<code>import net.dns</code>	DNS wire-format message parsing/building, name compression
<code>ip.sage</code>	<code>import net.ip</code>	IPv4 parsing/validation, CIDR subnets, private/loopback/multicast checks

9.12 Cryptography Libraries

SageLang ships with 6 cryptography modules in `lib/crypto/`:

Module	Import	Purpose
<code>hash.sage</code>	<code>import crypto.hash</code>	SHA-256, SHA-1, CRC-32 hash functions with hex output
<code>hmac.sage</code>	<code>import crypto.hmac</code>	HMAC (RFC 2104) with pluggable hash, constant-time comparison
<code>encoding.sage</code>	<code>import crypto.encoding</code>	Base64 (standard + URL-safe), hex encoding/decoding
<code>cipher.sage</code>	<code>import crypto.cipher</code>	XOR cipher, RC4 stream cipher, PKCS#7 padding, CBC/CTR mode helpers
<code>rand.sage</code>	<code>import crypto.rand</code>	xoshiro256** PRNG, UUID v4, random bytes/strings/hex, shuffle
<code>password.sage</code>	<code>import crypto.password</code>	PBKDF2-HMAC key derivation, password hashing/verification

9.13 Machine Learning Libraries

SageLang ships with 10 PyTorch-style machine learning modules in `lib/ml/`:

Module	Import	Purpose
<code>tensor.sage</code>	<code>import ml.tensor</code>	N-dimensional tensors, element-wise ops, matmul, reductions, activations
<code>nn.sage</code>	<code>import ml.nn</code>	Neural network layers (Linear, ReLU, Sigmoid, Dropout), Sequential model
<code>optim.sage</code>	<code>import ml.optim</code>	SGD (momentum), Adam optimizer, learning rate schedulers
<code>loss.sage</code>	<code>import ml.loss</code>	MSE, cross-entropy, Huber, L1, hinge, KL divergence
<code>data.sage</code>	<code>import ml.data</code>	Dataset/DataLoader, batching, normalization, train/test split
<code>debug.sage</code>	<code>import ml.debug</code>	Weight stats, histograms, activation analysis, gradient checking, training diagnostics
<code>viz.sage</code>	<code>import ml.viz</code>	SVG chart generation (loss curves, weight distributions, attention heatmaps, architecture diagrams)
<code>monitor.sage</code>	<code>import ml.monitor</code>	Live training monitor, progress bars, memory snapshots, throughput, checkpoints
<code>gpu_accel.sage</code>	<code>import ml.gpu_accel</code>	GPU offload for tensor ops and training; bridges ML workloads to the Vulkan/OpenGL GPU backend
<code>npu.sage</code>	<code>import ml.npu</code>	NPU backend for Qualcomm Hexagon (SNPE), Samsung Exynos (ONE), NNAPI, and ARM NEON SIMD fallback

Native ML Backend (`ml_native`) `ml_native` is a built-in native C module (not a `lib/ml/` Sage file) that exposes 24 high-performance functions directly from `src/c/ml_backend.c`:

`matmul`, `add`, `scale`, `relu`, `gelu`, `silu`, `sigmoid`, `softmax`, `layer_norm`, `rms_norm`, `cross_entropy`, `adam_update`, `clip_grad`, `benchmark`, `train_step`, `forward_pass`, `load_weights`, `gpu_available`, `set_gpu_threshold`, `cpu_count`, `auto_parallel`, `set_threads`, `set_parallel_threshold`, `get_threads`

Backpropagation and training:

- `ml_native.train_step()` — C-level combined forward pass + backward pass + SGD weight update for transformer training. Avoids Python/Sage overhead in the hot loop.
- `ml_native.forward_pass()` — Runs inference using the same computation graph as `train_step`, ensuring numerical consistency between training and evaluation.
- `ml_native.load_weights(path)` — Loads trained weights from a binary file via native C parser; no Sage I/O overhead.

Standalone C trainer: `make train-c` builds a standalone C training binary (no Sage runtime required) that uses `ml_backend.c` directly. Auto-detects cuBLAS GPU acceleration and ARM NEON SIMD; also works on mobile via Termux + proot ARM64. Additional targets: `make train-sage` (Sage interpreter), `make chatbot-c` / `make chatbot-llvm` (compile chatbots), `make sl-tq-chat` (SL-TQ-LLM generative chatbot).

9.14 CUDA Libraries

SageLang ships with 4 CUDA abstraction modules in `lib/cuda/`:

Module	Import	Purpose
<code>device.sage</code>	<code>import cuda.device</code>	GPU device descriptors, compute capability, feature detection
<code>memory.sage</code>	<code>import cuda.memory</code>	GPU memory allocation, typed tensors, memory pools
<code>kernel.sage</code>	<code>import cuda.kernel</code>	Kernel definition, launch parameters, occupancy analysis
<code>stream.sage</code>	<code>import cuda.stream</code>	CUDA streams, events, multi-stream execution plans

9.15 Standard Library (`lib/std/`)

SageLang ships with 23 general-purpose standard library modules in `lib/std/`:

Core utilities:

Module	Import	Purpose
<code>regex.sage</code>	<code>import std.regex</code>	Regular expression engine (<code>.</code> , <code>*</code> , <code>+</code> , <code>?</code> , <code>[]</code> , <code>\\d</code> , <code>\\w</code> , <code>\\s</code>)
<code>datetime.sage</code>	<code>import std.datetime</code>	Date/time creation, ISO 8601, arithmetic, comparison
<code>log.sage</code>	<code>import std.log</code>	Structured logging (TRACE-FATAL), handlers, child loggers
<code>argparse.sage</code>	<code>import std.argparse</code>	CLI argument parser (flags, options, positionals)
<code>compress.sage</code>	<code>import std.compress</code>	RLE, LZ77, delta encoding/decoding
<code>process.sage</code>	<code>import std.process</code>	Environment, path manipulation, exit codes, timers
<code>unicode.sage</code>	<code>import std.unicode</code>	UTF-8, character classification, case conversion, trim
<code>fmt.sage</code>	<code>import std.fmt</code>	Number/string formatting, templates, table output
<code>testing.sage</code>	<code>import std.testing</code>	Test runner, assertions, benchmarks
<code>db.sage</code>	<code>import std.db</code>	In-memory DB (CRUD, joins, aggregation, pagination)
<code>signal.sage</code>	<code>import std.signal</code>	Event bus (pub/sub), atexit handlers

Type system extensions:

Module	Import	Purpose
<code>enum.sage</code>	<code>import std.enum</code>	Enums, tagged unions, Result (Ok/Err), Option (Some/None)
<code>trait.sage</code>	<code>import std.trait</code>	Interface/trait system, behavioral contracts

Concurrency primitives:

Module	Import	Purpose
<code>channel.sage</code>	<code>import std.channel</code>	Go-style channels, send/recv/select, fan-in/out
<code>threadpool.sage</code>	<code>import std.threadpool</code>	Work queue, parallel map, futures/promises
<code>atomic.sage</code>	<code>import std.atomic</code>	Atomic integers, CAS, spin locks, counters
<code>rwlock.sage</code>	<code>import std.rwlock</code>	Read-write locks, scoped lock helpers
<code>condvar.sage</code>	<code>import std.condvar</code>	Condition variables, barriers, latches, semaphores

Developer tooling:

Module	Import	Purpose
<code>debug.sage</code>	<code>import std.debug</code>	Value inspection, trace logging, watch expressions
<code>profiler.sage</code>	<code>import std.profiler</code>	Hierarchical timing, hotspots, benchmark runner
<code>docgen.sage</code>	<code>import std.docgen</code>	Doc extraction from comments, markdown generation
<code>build.sage</code>	<code>import std.build</code>	Project config, dependencies, semver, build targets
<code>interop.sage</code>	<code>import std.interop</code>	FFI helpers, C types, pack/unpack, platform detection

9.16 LLM / Neural Network Libraries

SageLang ships with 19 LLM/neural network modules in `lib/llm/` for building and training language models:

Backpropagation note: For performance-critical training loops, use `ml_native.train_step()` (C-level forward+backward+SGD), `ml_native.forward_pass()` (inference, same computation graph), and `ml_native.load_weights(path)` (native weight loading) instead of the pure-Sage `llm.train` module. See Section 9.13 Native ML Backend.

Module	Import	Purpose
<code>config.sage</code>	<code>import llm.config</code>	Model configs (tiny to Llama-13B), param counting, memory estimation
<code>tokenizer.sage</code>	<code>import llm.tokenizer</code>	Character, word, and BPE tokenizers
<code>embedding.sage</code>	<code>import llm.embedding</code>	Token embeddings, sinusoidal/learned/RoPE positional encodings
<code>attention.sage</code> <code>transformer.sage</code>	<code>import llm.attention</code> <code>import llm.transformer</code>	Multi-head self-attention, KV cache Transformer blocks, LayerNorm/RMSNorm, FFN, model assembly
<code>generate.sage</code>	<code>import llm.generate</code>	Text generation (greedy, top-k, top-p, beam search)
<code>train.sage</code>	<code>import llm.train</code>	Training loops, LR schedules, cross-entropy, perplexity

Module	Import	Purpose
<code>agent.sage</code>	<code>import llm.agent</code>	Agentic framework (tools, CoT, memory, planning, multi-agent)
<code>prompt.sage</code>	<code>import llm.prompt</code>	Chat formatting (ChatML/Llama/Alpaca), templates, few-shot
<code>lora.sage</code>	<code>import llm.lora</code>	LoRA fine-tuning adapters, merge-back
<code>quantize.sage</code>	<code>import llm.quantize</code>	Int8/int4 quantization, error analysis
<code>engram.sage</code>	<code>import llm.engram</code>	Persistent neural memory (working/episodic/semantic/procedural)
<code>rag.sage</code>	<code>import llm.rag</code>	Document chunking, keyword retrieval, context assembly, summarization
<code>dpo.sage</code>	<code>import llm.dpo</code>	Direct Preference Optimization, ORPO, preference pairs, reward models
<code>gguf.sage</code>	<code>import llm.gguf</code>	GGUF v3 export for Ollama/llama.cpp, Modelfile gen, quantization
<code>gguf_import.sage</code>	<code>import llm.gguf_import</code>	Import GGUF models from Ollama into Sage; converts weights to native tensor format
<code>turboquant.sage</code>	<code>import llm.turboquant</code>	TurboQuant near-optimal vector quantization (ICLR 2026): two-stage PolarQuant (random rotation + MSE-optimal scalar quantization) + QJL (1-bit residual correction); KV cache compression at 3-bit with 6x memory reduction
<code>autoresearch.sage</code>	<code>import llm.autoresearch</code>	Karpathy-style autonomous research agent; ratchet loop (propose → train → evaluate → accept/reject); built-in scale/choice/perturb strategies; research journal; multi-agent session merging
<code>evolve.sage</code>	<code>import llm.evolve</code>	Self-evolving neural architecture; progressive growth from seed (64d/1L/98K params) to ancient (512d/8L/67M params); auto-plateau detection; weight padding for width growth; identity-init for depth growth

9.17 Agent Framework (lib/agent/)

Module	Import	Purpose
<code>core.sage</code>	<code>import agent.core</code>	ReAct agent loop, tool dispatch, scratchpad, prompt building
<code>tools.sage</code>	<code>import agent.tools</code>	Pre-built tools (file I/O, code analysis, search, system)

Module	Import	Purpose
<code>planner.sage</code>	<code>import agent.planner</code>	Task decomposition with dependency DAG, auto-execution
<code>router.sage</code>	<code>import agent.router</code>	Multi-agent orchestrator, capability routing, pipelines
<code>supervisor.sage</code>	<code>import agent.supervisor</code>	Supervisor-Worker control plane, workflow engine, retries
<code>critic.sage</code>	<code>import agent.critic</code>	Verification loops, rule/LLM validators, composite checks
<code>schema.sage</code>	<code>import agent.schema</code>	Typed tool interfaces, parameter validation, bounded execution
<code>trace.sage</code>	<code>import agent.trace</code>	SFT trace recording, training data generation
<code>grammar.sage</code>	<code>import agent.grammar</code>	Grammar-constrained decoding, tool call/JSON/Sage validation
<code>sandbox.sage</code>	<code>import agent.sandbox</code>	Program-aided reasoning, code execution, math eval
<code>tot.sage</code>	<code>import agent.tot</code>	Tree of Thoughts, BFS/best-first search, state rollbacks
<code>semantic_router.sage</code>	<code>import agent.semantic_router</code>	Fast command dispatch, keyword routing, LLM bypass

9.18 Chatbot Framework (`lib/chat/`)

Module	Import	Purpose
<code>bot.sage</code>	<code>import chat.bot</code>	Conversation management, intents, middleware, LLM responses
<code>session.sage</code>	<code>import chat.session</code>	Multi-session store, history, export (text/JSON)
<code>persona.sage</code>	<code>import chat.persona</code>	Pre-built personas (SageDev, Teacher, Debugger, etc.)

9.19 UI Widget Library (`lib/graphics/ui.sage`)

Immediate-mode GPU UI system for building application interfaces:

```
import gpu
import graphics.ui

gpu.init_windowed("App", 800, 600, "My App", false)
let ctx = ui.ui_create()

while not gpu.window_should_close():
    gpu.poll_events()
    ui.ui_begin_frame(ctx)

    ui.ui_panel(ctx, 10, 10, 300, 400, "Settings")
    if ui.ui_button(ctx, 20, 50, 120, 30, "Apply"):
        print "Applied!"

    let volume = ui.ui_slider(ctx, 20, 100, 260, "Volume", 0.75)
    let dark = ui.ui_checkbox(ctx, 20, 140, "Dark Mode", true)
```

```
ui.ui_end_frame(ctx)
```

Available widgets:

Widget	Function	Returns
Label	<code>ui_label(ctx, x, y, text)</code>	nil
Button	<code>ui_button(ctx, x, y, w, h, label)</code>	bool (clicked)
Panel	<code>ui_panel(ctx, x, y, w, h, title)</code>	nil
Window	<code>ui_window(ctx, x, y, w, h, title)</code>	dict (content area)
Checkbox	<code>ui_checkbox(ctx, x, y, label, checked)</code>	bool (new state)
Slider	<code>ui_slider(ctx, x, y, w, label, value)</code>	number (0.0-1.0)
Scrollbar	<code>ui_scrollbar_v(ctx, x, y, h, content_h, scroll)</code>	number (0.0-1.0)
Menu	<code>ui_menu_button(ctx, x, y, w, h, label, items)</code>	int (item index or -1)
Text Input	<code>ui_text_input(ctx, x, y, w, label, text)</code>	string
Progress	<code>ui_progress(ctx, x, y, w, h, value, label)</code>	nil
Separator	<code>ui_separator(ctx, x, y, w)</code>	nil
Tooltip	<code>ui_tooltip(ctx, text)</code>	nil

Theming: `ctx["theme"]` is a dict with 20+ configurable colors (`bg`, `accent`, `text`, `border`, etc.) and sizes (`padding`, `font_size`, `title_height`, etc.). Call `ui.ui_default_theme()` for defaults.

9.20 AVR Assembler & Arduino Uno Support (core/boards/AVR/)

SageLang v4.1.7 introduces a complete two-pass assembler package targeting the **AVR 8-bit RISC architecture (ATmega328P / ATmega328PB)** such as the Arduino Uno R3.

This includes: - **avr_assembler**: Parses AVR assembly, handles label resolution, relative offsets for branching (`brne`, `rjmp`, `rcall`), and resolves preprocessor directives (`.org`, `.equ`, `.byte`). - **avr_opcodes**: Encoder translating mnemonics to exact 16-bit instructions matching the Microchip datasheet. - **avr_hex**: Intel HEX (I8HEX) emitter generating `.hex` files ready for `avrdude` flashing.

Standard I/O registers resolved automatically:

- `DDRB`, `PORTB`, `PINB`, `DDRC`, `PORTC`, `PINC`, `DDRD`, `PORTD`, `PIND`, `SPL`, `SPH`, `SREG`

Example Usage:

```
import avr_assembler
import avr_hex
import io
```

```
let src = "
    .org 0x0000
    ldi r16, 0x20
```

```

    out DDRB, r16
    out PORTB, r16
    ret
"
let words = avr_assembler.assemble(src)
let hex_txt = avr_hex.emit_hex(words, 0)
print(hex_txt)

```

9.21 ESP32 Board Support (core/boards/ESP32/, core/lib/esp32.sage)

SageLang v4.2.12 supports the **classic ESP32** (verified on ESP32-D0WD-V3, ESP-WROOM-32 / DevKitC, 4MB flash): dual-core Xtensa LX6 @ 240MHz with 2.4GHz-only WiFi.

This includes:

- **core/lib/esp32.sage**: Board support module with chip constants, GPIO validation rules (pads 0..39 minus absent 20/24/28-31, input-only 34..39, strapping 0/2/5/12/15, flash-bound 6/7/8/11), UART0 TX/RX defaults, standard flash partition offsets, and esptool flashing recipe parameters.
- **core/boards/ESP32/**: Board package with `__init__.sage`, host-runnable `test_smoke.sage`, and bring-up firmware sources (`examples/hello.sage`, `examples/blink.sage`, `examples/adc_read.sage`, `examples/sysinfo.sage`, `examples/deep_sleep.sage`).

Exported Hardware Constants & Predicates:

- **Identity**: `CHIP_NAME` ("ESP32"), `CHIP_VARIANT` ("ESP32-D0WD-V3"), `CPU_ARCH` ("Xtensa LX6"), `CPU_CORES` (2), `CPU_FREQ_MHZ` (240), `HAS_LP_CORE` (true)
- **Memory Map**: Internal SRAM accommodates 98 KB of `.bss` residing safely at `0x3ffce000-0x3ffe6938` within the `0x40000000` limit. The `.data` copy during boot is strictly required to prevent reset loops, even if landing in `IRAM0`.
- **Peripherals**: UART0 FIFO is located at `0x3FF40000` and UART0 STATUS is at `0x3FF4001C`. Scratch probing the STATUS register instead of FIFO can cause artificial reset loops. Note: When investigating hangs during OS bring-up, volatile marker stores may be reordered by the compiler, invalidating localization theories.
- **Radio**: `WIFI_BAND_GHZ` (2.4), `WIFI_SUPPORTS_5GHZ` (false), `HAS_BT` (true), `HAS_BLE` (true)
- **SPI Flash Layout**: `FLASH_SIZE_BYTES` (4194304), `FLASH_BOOTLOADER_OFFSET` (4096), `FLASH_PART_TABLE_OFFSET` (32768), `FLASH_APP_OFFSET` (65536)
- **GPIO Model**: `GPIO_MIN` (0), `GPIO_MAX` (39), `GPIO_ABSENT` ([20, 24, 28, 29, 30, 31]), `GPIO_INPUT_ONLY_FIRST` (34), `GPIO_STRAPPING` ([0, 2, 5, 12, 15]), `GPIO_FLASH_PINS` ([6, 7, 8, 11]), `GPIO_ADC1_PINS` ([32..39]), `GPIO_DAC_PINS` ([25, 26]), `GPIO_RTC_PADS` ([0, 2, 4, 12, 13, 14, 15, 25, 26, 27, 32..39])
- **Peripherals & Predicates**:
 - `pin_valid(pin)`: True for physical ESP32 pads.
 - `pin_can_output(pin)`: True for output-capable pads (excludes input-only pads 34..39 and absent pads).
 - `pin_is_input_only(pin)`: True for input-only pads (34..39).
 - `pin_is_strapping(pin)`: True for boot strapping pads (0, 2, 5, 12, 15).
 - `pin_is_flash(pin)`: True for flash-bound SPI pins (6, 7, 8, 11).
 - `adc1_channel(pin)`: Maps pad 32..39 to ADC1 channel 0..7 (-1 if unsupported; ADC1 avoids WiFi ADC2 conflict).
 - `adc_fullscale_mv(atten_db)`: Approximate full scale voltage in mV for 0, 2.5, 6, or 11 dB attenuation.
 - `adc_to_mv(raw, atten_db)`: Converts raw 12-bit ADC count (0..4095) to millivolts.
 - `touch_channel(pin)`: Maps pad to touch channel T0..T9 (-1 if unsupported).
 - `rtc_capable(pin)`: True for deep-sleep wake RTC pads.
 - `pwm_duty(percent, bits)`: Calculates floored duty count for LEDC PWM given percentage (0..100) and resolution bits.
 - `uart_needs_remap(uart)`: Returns true if UART default pins collide with flash (UART1 default TX 10 / RX 9).

- `esptool_write_cmd(port, image)`: Generates verified `esptool` flash command string.
- `describe()`: One-line summary string for the board configuration.

Example Usage:

```
import esp32

print esp32.describe()
if esp32.pin_can_output(2):
    print "GPIO2 ready for LED drive"

let flash_app = esp32.flash_offset("app")
print "App partition offset: " + str(flash_app)

let ch = esp32.adc1_channel(34)
let mv = esp32.adc_to_mv(2048, 11)
print "ADC1 Channel: " + str(ch) + " Voltage (mV): " + str(mv)

let duty = esp32.pwm_duty(50, 8)
print "50% Duty Count (8-bit): " + str(duty)

let cmd = esp32.esptool_write_cmd("/dev/ttyUSB0", "firmware.bin")
print "Flash Command: " + cmd
```

Part 13: Self-Hosting (Phase 13)

SageLang is **self-hosted**: the lexer, parser, and interpreter have been ported from C to Sage itself. The self-hosted pipeline can execute `.sage` programs using only the C interpreter as a bootstrap host.

13.1 Architecture

The self-hosted interpreter lives in `src/sage/` and consists of:

File	Description	Size
<code>token.sage</code>	Token type constants	~50 lines
<code>lexer.sage</code>	Tokenizer with indentation tracking	~300 lines
<code>ast.sage</code>	AST node constructors (dict-based)	~100 lines
<code>parser.sage</code>	Recursive descent parser	~700 lines
<code>interpreter.sage</code>	Tree-walking evaluator with module imports	~1050 lines
<code>errors.sage</code>	Rich error reporting (Rust/Elm-style)	~200 lines
<code>environment.sage</code>	Dict-based scope/environment	~35 lines
<code>sage.sage</code>	Full CLI entry point	~200 lines

13.2 Running Self-Hosted Code

```
cd src/sage
../../sage sage.sage program.sage
```

The bootstrap reads a `.sage` file, tokenizes it, parses it to an AST, and evaluates it — all in Sage running on the C interpreter.

13.3 Key Design Decisions

Dict-based value representation: While SageLang recently added native C-like enums (Phase 1.7), the self-hosted interpreter relies on dictionary-based representations for its AST nodes because it was designed before these features existed and for simplicity. All AST nodes, functions, classes, and instances are represented as dicts with an `__interp_type` field:

```
# A function value
let fn = {}
fn["__interp_type"] = "function"
fn["name"] = "add"
fn["params"] = ["a", "b"]
fn["body"] = body_ast
```

Control flow signals: Return, break, and continue are implemented as dict values:

```
# return 42 → {"kind": "return", "value": 42}
# break → {"kind": "break"}
# continue → {"kind": "continue"}
# normal → {"kind": "normal"}
```

GC must be disabled: The self-hosted interpreter creates many dict allocations, which can trigger GC segfaults. Always start with `gc_disable()`.

13.4 Feature Coverage

The self-hosted interpreter supports (~70% feature parity with C):

Fully implemented: bare `end` statements (optional block terminators), anonymous procedure expressions (`proc(params...): body [end]`), full CLI flag surface parity with C `main.c` (`-c`, `--jit`, `--aot`, `--sgvm`, `--lsp`), arithmetic, variables, control flow (`if/elif/else`, `while`, `for`), functions with closures, recursion, classes with inheritance, arrays, dicts, tuples, strings, slicing, indexing, `try/catch/finally`, `break/continue`, `raise`, module imports (`import X`, `import X as Y`, `from X import a, b`), property access/set, all bitwise operators (`&` | `^` | `~` | `<<` | `>>`), `match/case/default` pattern matching, `defer` statements (LIFO cleanup on scope exit), generators/yield (eager collection via `next()`), 35+ builtins.

For self-hosted LLVM codegen specifically (`src/sage/llvm_backend.sage`), `from X import Y` constant imports are resolved during code generation for foldable top-level `let` values (including `as` aliases), matching the C LLVM backend behavior for cross-module constants.

Data Representation: The self-hosted interpreter relies on dictionary-based representations for its AST nodes, functions, classes, and instances, tagged with an `__interp_type` field, instead of native enums.

Delegated to host runtime: GC control (`gc_collect`, `gc_enable`, `gc_disable`, `gc_stats`), FFI (`ffi_open`, `ffi_close`, `ffi_call`, `ffi_sym`), memory access (`mem_alloc`, `mem_free`, `mem_read`, `mem_write`, `mem_size`, `addressof`), networking (via host `import` of native modules).

Stub/partial: async threading in the self-hosted interpreter evaluates synchronously (AST-backed path).

Not implemented: true coroutine-based generators (uses eager collection instead), GPU module (uses host `import gpu` directly).

Safety: while loop iteration limit (1M), recursion depth limit (50000 in the self-hosted interpreter, backed by the host's stack-proximity guard), rich error messages with source context via `errors.sage`.

13.5 Test Suite

The self-hosted tree includes core parser/interpreter suites plus additional tooling, optimizer, backend, compiler, error-reporting, LSP, and CLI tests. The foundational bootstrap suites are:

Category	Tests	Coverage
test_lexer.sage	13	Token types, indentation, keywords
test_parser.sage	130	All AST node types, operator precedence
test_interpreter.sage	18	Evaluation, scoping, closures, classes
test_bootstrap.sage	18	End-to-end: source → tokens → AST → result

The full test suite (interpreter + compiler + self-hosted tooling) totals **2060+ tests**. Execute the unit tests using `bash testsuite/unit/run_tests.sh`.

Part 10: Native Standard Library Modules (Phase 11)

SageLang provides built-in native modules that are pre-loaded into the module cache and available via `import` without any file on disk.

10.1 Math Module

```
import math

print math.sqrt(16)      # 4
print math.sin(math.pi) # ~0
print math.floor(3.7)    # 3
print math.ceil(3.2)     # 4
print math.abs(-5)       # 5
print math.pow(2, 10)    # 1024
print math.log(math.e)   # 1
math.printm("10 + 20") # Show work for 10 + 20
```

Available functions: `sqrt`, `sin`, `cos`, `tan`, `floor`, `ceil`, `abs`, `pow`, `log`, `printm` Constants: `pi` (3.14159...), `e` (2.71828...)

10.2 IO Module

```
import io

io.writefile("test.txt", "Hello, World!")
let content = io.readfile("test.txt")
print content # Hello, World!

io.appendfile("test.txt", "\nLine 2")
print io.exists("test.txt") # true

io.rename("test.txt", "renamed.txt")
io.remove("renamed.txt")

# Binary I/O
let buf = bytes("Hello\n")
io.writebytes("data.bin", buf) # Write Bytes value (also accepts Array)
let r = io.readbytes("data.bin") # Read as Bytes value
io.appendbytes("log.bin", buf) # Append Bytes value
```

Available functions: `readfile`, `writefile`, `appendfile`, `exists`, `remove`, `rename`, `readbytes`, `writebytes`, `appendbytes`

Note: `io.readbytes` returns a `Bytes` value (byte buffer), and `io.writebytes`/`io.appendbytes` accept either `Bytes` or `Array` values.

10.3 String Module

```
import string

print string.char(65)          # A
print string.ord("A")          # 65
print string.startswith("hello", "he") # true
print string.endswith("hello", "lo")   # true
print string.contains("hello", "ell")   # true
print string.repeat("ab", 3)            # ababab
print string.reverse("hello")           # olleh
```

Available functions: `char`, `ord`, `startswith`, `endswith`, `contains`, `repeat`, `reverse`

10.4 Sys Module

```
import sys

print sys.platform() # linux, darwin, or windows
print sys.version()  # Sage version string
let path = sys.env("HOME")
print path
```

Available functions: `args`, `exit`, `platform`, `version`, `env`, `setenv`

10.5 FAT Module

The native `fat` module provides early FAT filesystem parsing helpers for image inspection and kernel/boot tooling.

```
import os.fat
import io

let boot = io.readbytes("disk.img")
let info = fat.parse_boot_sector(boot)

print info["fat_type"]          # FAT8 / FAT12 / FAT16 / FAT32
print info["cluster_count"]
print fat.cluster_to_lba(info, 2) # First data cluster -> LBA/sector index
```

Available functions: - `parse_boot_sector(bytes_array)` - `probe(path)` - `cluster_to_lba(info, cluster)` - `fat_entry_offset(info, cluster)`

Constants: - `FAT8`, `FAT12`, `FAT16`, `FAT32`

Current scope: - Boot sector parsing, FAT type detection, and layout math for FAT8/12/16/32. - Intended as a foundation for follow-up directory walking, FAT chain traversal, and read/write support.

Part 11: Concurrency and Async/Await (Phase 11)

11.1 Thread Module

The `thread` module provides low-level threading primitives backed by pthreads.

```
import thread

proc worker(name):
  print name
  return 42

# Spawn a thread
let t = thread.spawn(worker, "hello")

# Wait for result
let result = thread.join(t)
print result # 42
```

Thread Functions

Function	Description
<code>thread.spawn(proc, args...)</code>	Spawn a new thread running <code>proc</code> with given arguments
<code>thread.join(t)</code>	Block until thread <code>t</code> completes, return its result
<code>thread.mutex()</code>	Create a new mutex
<code>thread.lock(m)</code>	Acquire mutex <code>m</code> (blocks if held)
<code>thread.unlock(m)</code>	Release mutex <code>m</code>
<code>thread.sleep(ms)</code>	Sleep for <code>ms</code> milliseconds
<code>thread.id()</code>	Return current thread identifier

Mutex Example

```
import thread

let m = thread.mutex()

proc critical_section():
  thread.lock(m)
  # ... protected code ...
  thread.unlock(m)
```

11.2 Async/Await

The `async proc` keyword declares a procedure that runs asynchronously. Calling it spawns a background thread and returns a thread handle. Use `await` to retrieve the result.

```
async proc compute(x):
  return x * x

# Calling an async proc spawns a thread
let future = compute(5)

# await blocks until the thread completes
let result = await future
print result # 25
```

Parallel Execution

```
async proc slow_add(a, b):
  return a + b

# Launch two computations in parallel
let f1 = slow_add(1, 2)
let f2 = slow_add(3, 4)

# Await both results
let r1 = await f1
let r2 = await f2
print r1 + r2 # 10
```

How It Works

1. `async proc` is parsed as `STMT_ASYNC_PROC` and sets `is_async = 1` on the `FunctionValue`
2. When called, the interpreter pre-evaluates arguments and spawns a thread via `thread_spawn_native`
3. The call returns a `VAL_THREAD` value (a thread handle)
4. `await` on a `VAL_THREAD` calls `pthread_join` and returns the thread's result value

11.3 GC Thread Safety

The garbage collector is protected by a pthread mutex. All GC operations (allocation, collection, marking, sweeping) acquire the lock, ensuring safe concurrent allocation from multiple threads.



Part 12: Developer Tooling (Phase 12)

SageLang includes a complete set of developer tools for interactive development, code quality, and editor integration.

12.1 REPL (Read-Eval-Print Loop)

Launch the interactive REPL by running `sage` with no arguments or with the `--repl` flag:

```
sage
sage --repl
```

The REPL supports multi-line blocks (indented code is automatically continued), error recovery (errors are displayed without exiting), and built-in commands.

Built-in Commands:

Command	Description
<code>:help</code>	Show REPL usage information
<code>:quit</code> / <code>:exit</code>	Exit the REPL
<code>:vars [prefix]</code>	List current REPL bindings, optionally filtered by prefix
<code>:type <expr></code>	Evaluate an expression and print its runtime type and value
<code>:load <file></code>	Execute a Sage file inside the current REPL session
<code>:reset</code>	Reset the REPL environment and module cache
<code>:pwd</code>	Print the current working directory
<code>:cd <dir></code>	Change the current working directory
<code>:cat <file></code>	Print file contents. Uses <code>is_safe_path</code> (blocks single quotes, wraps in quotes).

Command	Description
<code>:ls [dir]</code>	List directory contents. Uses <code>is_safe_path</code> .
<code>:sh <cmd></code>	Execute a shell command. Uses <code>is_safe_command</code> (blocks metacharacters like <code>;</code> , <code>&</code> , <code> </code> , <code>></code>).
<code>:gc</code>	Run garbage collection and print GC statistics
<code>:stats</code>	Show GC stats, stack depth, and CPU time

Note on REPL Security: The `:cat` and `:ls` commands validate arguments using `is_safe_path`, which allows spaces but blocks single quotes to prevent shell injection, wrapping the final argument in single quotes securely. The `:sh` command uses `is_safe_command` which allows both spaces and single quotes for arguments, but explicitly blocks shell metacharacters (`;`, `&`, `|`, `>`, etc.) for consistency with `sys.exec` hardening.

Example Session:

```
sage> let x = 10
sage> proc double(n):
....>     return n * 2
sage> print double(x)
20
sage> :quit
```

12.2 Code Formatter

The formatter normalizes indentation, spacing, and blank lines for consistent code style.

```
sage fmt program.sage           # Format file in place (Note: may incorrectly format << and >> as < <
sage fmt --check program.sage   # Check formatting (exit code 1 if changes needed)
```

Example:

```
$ sage fmt --check messy.sage
messy.sage: needs formatting
$ sage fmt messy.sage
Formatted: messy.sage
$ sage fmt --check messy.sage
messy.sage: already formatted
```

12.3 Linter

The linter performs static analysis with 13 rules across three categories:

```
sage lint program.sage
```

Rule Categories:

- [W001]: Unused variable warning.
- [W002]: Shadowed variable warning.
- [W003]: Unreachable code warning (e.g., after return/break/continue).
- [W004]: Empty block warning (e.g., following a colon).
- [S003]: Missing docstring warning (requires `##` preceding top-level procedures).
- [S004]: Trailing semicolon warning.
- [S005]: Multiple statements on a single line.

Category	Rules	Description
Errors	E001-E003	Syntax and structural errors
Warnings	W001-W005	Potential bugs and bad practices

Category	Rules	Description
Style	S001-S005	Code style and naming conventions

Notable Rules: - [W001], [W002]: Issued for unused variables and variables that shadow existing declarations in scope. - [W003]: Warns about unreachable code following a `return`, `break`, or `continue` statement at the same indentation level. - [W004]: Warns about empty blocks (e.g., when a line ending in a colon `:` is followed by an empty or improperly indented section). - [S003]: Enforces documentation conventions; top-level `proc` declarations must be immediately preceded by a `##` style docstring (comment). Top-level procedure doc comments (`## ...`) must be placed before decorators like `@inline`. Placing them between `@inline` and `proc` triggers a parser error. - [S004]: Warns about trailing semicolons (not used in SageLang). - [S005]: Warns when multiple statements are on a single line separated by semicolons.

12.4 OIS Package Management

SageLang includes a package manager accessed via the `--ois` flag. Commands include: - `sage --ois install <pkg>` - `sage --ois --update` - `sage --ois --uninstall <pkg>`

Example Output:

```
program.sage:5: W001 unused variable 'temp'
program.sage:12: E003 line too long (exceeds 120 characters)
program.sage:20: E001 Inconsistent indentation
```

12.5 Syntax Highlighting

SageLang provides editor support via TextMate grammars:

- **TextMate grammar:** `editors/sage.tmLanguage.json` works with any TextMate-compatible editor (VSCode, Sublime Text, etc.)
- **VSCode extension:** `editors/vscode/` provides full language support including syntax highlighting and language configuration

Installing the VSCode Extension:

1. Copy or symlink `editors/vscode/` into `~/.vscode/extensions/sage-lang/`
2. Restart VSCode
3. Open any `.sage` file to see syntax highlighting

12.6 Language Server Protocol (LSP)

The LSP server provides IDE-like features for any editor that supports the Language Server Protocol.

```
sage --lsp           # Start LSP server via main binary
sage-lsp            # Standalone LSP server binary
```

Capabilities:

Feature	Description
Diagnostics	Real-time error and warning reporting on file save
Completion	Keyword and symbol completions as you type
Hover	Type information and documentation on hover
Formatting	Format-on-save via <code>textDocument/formatting</code>

Editor Configuration (VSCode `settings.json`):

```
{
  "sage.lsp.path": "/path/to/sage-lsp"
}
```

12.7 CommonJS to ESM Transpiler (cjs2esm)

SageLang includes a bundled **CommonJS (CJS) to ECMAScript Modules (ESM)** transpiler written in SageLang. It preserves inline CommonJS execution order by default and can rewrite safe leading requires to native ESM imports.

Commands:

```
cjs2esm convert <input-path> [options]
cjs2esm inspect <file-path>
cjs2esm check <project-path>
cjs2esm report <project-path>
```

Options for convert:

- `--out <file|dir>`: Output file or directory. Without this option, the converter writes a same-directory `.mjs` file. Files containing `require()` are rejected when relocated until output-path rebasing is implemented.
- `--target <node18|node20|node22|node24>`: Target Node.js baseline (default: `node20`).
- `--mode <strict|compat|discord>`: Transformation mode (default: `compat`).
 - `strict`: Accept only leading static imports; reject dynamic, order-sensitive, cache-dependent, and shim-dependent constructs.
 - `compat`: Automatically inject `createRequire` and fallback shims where needed.
 - `discord`: Discord.js-focused compatibility mode.
- `--rewrite-dynamic-imports`: Rewrites eligible dynamic requires through an `async import()` helper.
- `--source-maps`: Generates `.mjs.map` source maps alongside converted files.
- `--update-package-json`: Updates or injects `"type": "module"` in `package.json`.
- `--dry-run`: Runs analysis without writing output.

Example Programs

Hello World:

```
print "Hello, World"
```

Factorial:

```
proc factorial(n):
  if n <= 1:
    return 1
  return n * factorial(n - 1)
```

```
print factorial(5) # 120
```

Fibonacci:

```
proc fib(n):
  if n <= 1:
    return n
  return fib(n - 1) + fib(n - 2)
```

```
for i in range(0, 10):
  print fib(i)
```

Object-Oriented:

```
class Rectangle:
    proc init(width, height):
        self.width = width
        self.height = height

    proc area():
        return self.width * self.height

let rect = Rectangle(5, 3)
print rect.area() # 15
```

Part 14: Networking (Phase 14)

SageLang provides four native networking modules originally planned to be backed by libcurl and OpenSSL. These are implemented in C (`src/net.c`) and registered as native modules. **Note:** Currently, only the `socket` and `tcp` modules have partial functional implementations. The `http` and `ssl` modules are experimental stubs/mocks returning hardcoded values.

14.1 Socket Module

Low-level POSIX socket operations with constants for address families and socket types.

```
import socket

# TCP client
let sock = socket.create(socket.AF_INET, socket.SOCK_STREAM, 0)
socket.connect(sock, "example.com", 80)
socket.send(sock, "GET / HTTP/1.0" + chr(13) + chr(10) + chr(13) + chr(10))
let response = socket.recv(sock, 4096)
print response
socket.close(sock)

# DNS resolution
let ip = socket.resolve("example.com")
print ip # "4.1.3.34"

# UDP
let udp = socket.create(socket.AF_INET, socket.SOCK_DGRAM, 0)
socket.sendto(udp, "hello", "4.1.3.1", 9999)
socket.close(udp)
```

Constants: `AF_INET`, `AF_INET6`, `SOCK_STREAM`, `SOCK_DGRAM`, `SOCK_RAW`, `IPPROTO_TCP`, `IPPROTO_UDP`

Functions (15): `create`, `bind`, `listen`, `accept`, `connect`, `send`, `recv`, `sendto`, `recvfrom`, `close`, `setopt`, `poll`, `resolve`, `getpeername`, `nonblock`

14.2 TCP Module

High-level TCP with automatic buffering and convenience functions.

```
import tcp

# Client
let conn = tcp.connect("example.com", 80)
tcp.sendall(conn, "GET / HTTP/1.0" + chr(13) + chr(10) + chr(13) + chr(10))
```

```

let line = tcp.recvline(conn)
print line
tcp.close(conn)

```

```

# Server
let server = tcp.listen("4.1.3.0", 8080, 5)
let client = tcp.accept(server)
tcp.send(client, "Hello from Sage!")
tcp.close(client)
tcp.close(server)

```

Functions (9): connect, listen, accept, send, recv, sendall, recvall, recvline, close

14.3 HTTP Module

Status: Experimental Stubs/Mocks. The HTTP module is currently incomplete. Functions like `http.get` and `http.post` return hardcoded mock string responses rather than performing real network operations via `libcurl`.

HTTP/HTTPS client API (planned via `libcurl`). When fully implemented, request functions will return a dict with `status`, `body`, and `headers` keys. Currently, they return mock strings.

```

import http

# Simple GET
let resp = http.get("https://httpbin.org/get")
print resp    # "HTTP GET response"

# POST with options
let opts = {"timeout": 30, "headers": {"Content-Type": "application/json"}}
let resp2 = http.post("https://httpbin.org/post", "{}", opts)

# Download file
http.download("https://example.com/file.txt", "/tmp/file.txt")

# URL encoding
print http.escape("hello world")    # "hello%20world"
print http.unescape("hello%20world") # "hello world"

```

Functions (9): get, post, put, delete, patch, head, download, escape, unescape

Options dict keys: `timeout` (seconds), `follow` (bool, follow redirects), `verify` (bool, SSL verification), `user_agent` (string), `headers` (dict of header name→value pairs), `cainfo` (CA certificate path)

14.4 SSL Module

Status: Unimplemented Stub. The OpenSSL TLS/SSL bindings for encrypted socket communication are currently an empty module stub in `src/net.c`.

```

import socket
import ssl

let sock = socket.create(socket.AF_INET, socket.SOCK_STREAM, 0)
socket.connect(sock, "example.com", 443)

let ctx = ssl.context()
let s = ssl.wrap(ctx, sock)
ssl.connect(s)

```

```

ssl.send(s, "GET / HTTP/1.1" + chr(13) + chr(10) + "Host: example.com" + chr(13) + chr(10) + chr(13) + chr(10))
let data = ssl.recv(s, 4096)
print data

ssl.shutdown(s)
ssl.free(s)
ssl.free_context(ctx)
socket.close(sock)

```

Functions (13): context, load_cert, wrap, connect, accept, send, recv, shutdown, free, free_context, error, peer_cert, set_verify

14.5 Important Notes

- **No escape sequences:** Sage strings are raw. Use `chr(13) + chr(10)` for CRLF, `chr(34)` for double-quote, `chr(92)` for backslash.
- **HTTP/SSL Mocks:** As of this version, the `http` module returns mock responses and the `ssl` module is an empty stub. While the API is documented above, real `libcurl` and `OpenSSL` integrations are not fully functional in the native backend.
- **Build requirement:** `libcurl` and `openssl` development libraries must be installed for desktop builds to succeed, even though they are currently unused by the runtime stubs.

Part 15: JSON Library (cJSON Port)

SageLang includes a complete 1:1 port of Dave Gamble's cJSON library in `lib/json.sage` (~1,050 lines). It uses a linked-list tree structure mirroring the C original.

15.1 Basic Usage

```

gc_disable() # Required for class-heavy code

from json import cJSON_Parse, cJSON_Print, cJSON_PrintUnformatted
from json import cJSON_GetObjectItem, cJSON_GetStringValue, cJSON_GetNumberValue
from json import cJSON_ToSage, cJSON_Delete

# Parse JSON string (use chr(34) for double-quotes in Sage)
let json_str = "{" + chr(34) + "name" + chr(34) + ":" + chr(34) + "Alice" + chr(34) + "," + chr(34) + "age" + chr(34) + ": 30" + chr(34) + "}"
let root = cJSON_Parse(json_str)

# Query values
let name = cJSON_GetStringValue(cJSON_GetObjectItem(root, "name"))
let age = cJSON_GetNumberValue(cJSON_GetObjectItem(root, "age"))
print name # Alice
print age  # 30

# Print JSON
print cJSON_Print(root) # Formatted (pretty-printed)
print cJSON_PrintUnformatted(root) # Compact

# Convert to native Sage dict
let native = cJSON_ToSage(root)
print native["name"] # Alice

cJSON_Delete(root) # No-op in GC language, included for API compatibility

```

15.2 Creating JSON

```
gc_disable()
from json import cJSON_CreateObject, cJSON_CreateArray
from json import cJSON_AddStringToObject, cJSON_AddNumberToObject
from json import cJSON_AddBoolToObject, cJSON_AddNullToObject
from json import cJSON_AddArrayToObject, cJSON_AddItemToArray
from json import cJSON_CreateNumber, cJSON_PrintUnformatted

let obj = cJSON_CreateObject()
cJSON_AddStringToObject(obj, "name", "Bob")
cJSON_AddNumberToObject(obj, "age", 25)
cJSON_AddBoolToObject(obj, "active", true)
cJSON_AddNullToObject(obj, "deleted_at")

let tags = cJSON_AddArrayToObject(obj, "tags")
cJSON_AddItemToArray(tags, cJSON_CreateNumber(1))
cJSON_AddItemToArray(tags, cJSON_CreateNumber(2))

print cJSON_PrintUnformatted(obj)
# {"name":"Bob","age":25,"active":true,"deleted_at":null,"tags":[1,2]}
```

15.3 Sage-Native Conversion

```
gc_disable()
from json import cJSON_FromSage, cJSON_ToSage, cJSON_Print

# Native Sage value → cJSON tree
let data = {"users": [{"name": "Alice"}, {"name": "Bob"}], "count": 2}
let tree = cJSON_FromSage(data)
print cJSON_Print(tree)

# cJSON tree → native Sage value
let back = cJSON_ToSage(tree)
print back["users"][0]["name"] # Alice
```

15.4 API Reference

Category	Functions
Parsing	cJSON_Parse, cJSON_ParseWithLength, cJSON_GetErrorPtr
Printing	cJSON_Print, cJSON_PrintUnformatted, cJSON_PrintBuffered
Creation	cJSON_CreateNull/True/False/Bool/Number/String/Raw/Array/Obj CreateIntArray/DoubleArray/FloatArray/StringArray
Query	cJSON_GetArraySize, GetArrayItem, GetObjectItem, GetObjectItemCaseSensitive, HasObjectItem, GetStringValue, GetNumberValue
Type Checks	cJSON_IsInvalid/False/True/Bool/Null/Number/String/Array/Obj
Array Ops	AddItemToArray, InsertItemInArray, DetachItemFromArray, DeleteItemFromArray, ReplaceItemInArray

Category	Functions
Object Ops	AddItemToObject/CS, DetachItemFromObject/CaseSensitive, DeleteItemFromObject/CaseSensitive, ReplaceItemInObject/CaseSensitive
Helpers	cJSON_AddNullToObject, AddTrue/False/Bool/Number/String/Raw/Array/ObjectToObject
Utility	cJSON_Duplicate, Compare, Minify, Delete, SetValuestring, SetNumberHelper, Version
Sage Extras	cJSON_ToSage (tree→native), cJSON_FromSage (native→tree)

15.5 Important Notes

- **GC must be disabled** (`gc_disable()`) when using `json.sage` — class-heavy code triggers GC seg-faults.
- `cJSON_GetObjectItem` does **case-insensitive** matching (uses `lower()`). Use `cJSON_GetObjectItemCaseSensitive` for exact matching.
- `cJSON_Delete` is a no-op (Sage has GC), included for API compatibility with C cJSON.
- **Type constants:** `cJSON_Invalid=0`, `cJSON_False=1`, `cJSON_True=2`, `cJSON_NULL=4`, `cJSON_Number=8`, `cJSON_String=16`, `cJSON_Array=32`, `cJSON_Object=64`, `cJSON_Raw=128`

Conclusion

SageLang is a comprehensive, well-structured scripting language for systems and embedded programming. Its design combines the approachability of Python with low-level control through direct memory access, FFI, and inline assembly. The phased development approach (Phases 1–18) progresses from core features through advanced topics including OOP, generators, compilation backends (C, LLVM IR, native assembly), concurrency, networking, JSON processing, and a self-hosted interpreter — all fully implemented and integrated.

Key Takeaways: - **Lexer + Parser + Interpreter** pipeline is modular and extensible. - **Value system** supports dynamic typing with heap-allocated complex objects and GC. - **Scoped environments** enable closures and lexical scoping. - **Exception handling, generators, and `async/await`** provide modern control flow. - **Module system** enables code reuse with native (C) and Sage library modules. - **Mark-and-sweep GC** manages memory automatically with thread safety. - **Multiple backends:** tree-walking interpreter, C codegen, LLVM IR, native assembly (x86-64, aarch64, rv64, mips). - **Networking:** POSIX sockets, TCP, HTTP/HTTPS (libcurl), SSL/TLS (OpenSSL). - **Self-hosted:** Lexer, parser, and interpreter ported to Sage with full bootstrap. - **Test suite:** interpreter/compiler coverage, JSON coverage, and broad self-hosted suites spanning parsing, tooling, optimization, codegen, compiler, LSP, and CLI behavior.

SageLang offers a practical balance between ease of use and systems-level control, making it ideal for prototyping, education, embedded scripting, and learning language implementation.

GPU Graphics (Vulkan + OpenGL)

SageLang includes a professional-grade graphics engine supporting both Vulkan and OpenGL 4.5 backends. The library is built in four layers:

1. **Pure C GPU API** (`gpu_api.h/gpu_api.c`) – Backend-agnostic ~100 functions, no interpreter dependency
2. **C native module** (`import gpu`) – Interpreter wrappers for Value-based argument handling
3. **Sage builder libraries** (`lib/vulkan.sage`, `lib/opengl.sage`, `lib/gpu.sage`) – Ergonomic helpers

4. Engine + UI libraries (lib/renderer.sage, lib/pbr.sage, lib/ui.sage, etc.) – Application-level systems

Three execution paths share the same GPU API:

- **Interpreter:** graphics.c wraps sgpu_* with Value extraction
- **LLVM compiled:** llvm_runtime.c provides 103 sage_rt_gpu_* bridge functions
- **Bytecode VM:** 30 dedicated BC_OP_GPU_* opcodes for frame-loop hot paths

Build Requirements

Vulkan and OpenGL support are auto-detected at build time via pkg-config. GLFW is required for windowed rendering:

```
# Auto-detect Vulkan + OpenGL (default)
make

# Force enable/disable
make VULKAN=1      # Force Vulkan
make VULKAN=0      # Disable Vulkan
make OPENGGL=1     # Force OpenGL
make OPENGGL=0     # Disable OpenGL
```

Without the Vulkan SDK, the gpu module loads in stub mode – all constants are available, functions return errors gracefully. Use import graphics.opengl instead of import gpu for the OpenGL 4.5 backend (same API, different initialization).

Quick Start: Empty Window

```
import gpu

gpu.init_windowed("My App", 800, 600, "Window Title", false)
print gpu.device_name()

let attach = {}
attach["format"] = gpu.FORMAT_SWAPCHAIN
attach["load_op"] = gpu.LOAD_CLEAR
attach["store_op"] = gpu.STORE_STORE
attach["initial_layout"] = gpu.LAYOUT_UNDEFINED
attach["final_layout"] = gpu.LAYOUT_PRESENT
let rp = gpu.create_render_pass([attach])
let framebuffers = gpu.create_swapchain_framebuffers(rp)

let cmd_pool = gpu.create_command_pool()
let cmd = gpu.create_command_buffer(cmd_pool)
let img_sem = gpu.create_semaphore()
let rdr_sem = gpu.create_semaphore()
let fence = gpu.create_fence(true)

while gpu.window_should_close() == false:
    gpu.poll_events()
    gpu.wait_fence(fence)
    gpu.reset_fence(fence)
    let idx = gpu.acquire_next_image(img_sem)
    gpu.begin_commands(cmd)
    gpu.cmd_begin_render_pass(cmd, rp, framebuffers[idx], [[0.1, 0.1, 0.2, 1.0]])
    gpu.cmd_end_render_pass(cmd)
```

```

    gpu.end_commands(cmd)
    gpu.submit_with_sync(cmd, img_sem, rdr_sem, fence)
    gpu.present(idx, rdr_sem)

gpu.device_wait_idle()
gpu.shutdown_windowed()

Quick Start: Triangle

import gpu

gpu.init_windowed("Triangle", 800, 600, "Sage Triangle", false)

# Load compiled SPIR-V shaders
let vert = gpu.load_shader("triangle.vert.spv", gpu.STAGE_VERTEX)
let frag = gpu.load_shader("triangle.frag.spv", gpu.STAGE_FRAGMENT)

# Render pass with swapchain format
let attach = {}
attach["format"] = gpu.FORMAT_SWAPCHAIN
attach["load_op"] = gpu.LOAD_CLEAR
attach["store_op"] = gpu.STORE_STORE
attach["initial_layout"] = gpu.LAYOUT_UNDEFINED
attach["final_layout"] = gpu.LAYOUT_PRESENT
let rp = gpu.create_render_pass([attach])

# Graphics pipeline
let layout = gpu.create_pipeline_layout([], 0)
let cfg = {}
cfg["layout"] = layout
cfg["render_pass"] = rp
cfg["vertex_shader"] = vert
cfg["fragment_shader"] = frag
cfg["topology"] = gpu.TOPO_TRIANGLE_LIST
cfg["cull_mode"] = gpu.CULL_NONE
let pipeline = gpu.create_graphics_pipeline(cfg)

# Render loop draws 3 hardcoded vertices
# gpu.cmd_bind_graphics_pipeline(cmd, pipeline)
# gpu.cmd_draw(cmd, 3, 1, 0, 0)

```

GPU Module API Reference

Context Lifecycle

Function	Description
<code>gpu.has_vulkan()</code>	Returns true if Vulkan is available
<code>gpu.has_window</code>	true if GLFW windowed mode is available
<code>gpu.initialize(name, validation?)</code>	Initialize headless Vulkan context
<code>gpu.init_windowed(name, w, h, title, validation?)</code>	Initialize windowed Vulkan + GLFW + swapchain
<code>gpu.shutdown()</code>	Destroy headless context
<code>gpu.shutdown_windowed()</code>	Destroy window + context
<code>gpu.device_name()</code>	GPU device name string

Function	Description
<code>gpu.device_limits()</code>	Dict of device limits

Windowing

Function	Description
<code>gpu.window_should_close()</code>	Check if window close requested
<code>gpu.poll_events()</code>	Process window events
<code>gpu.swapchain_extent()</code>	Dict {width, height}
<code>gpu.swapchain_image_count()</code>	Number of swapchain images
<code>gpu.acquire_next_image(semaphore)</code>	Get next image index (-1 on error)
<code>gpu.present(image_index, wait_semaphore)</code>	Present rendered image
<code>gpu.create_swapchain_framebuffers(render_passes)</code>	Array of framebuffer handles

Resources

Function	Description
<code>gpu.create_buffer(size, usage, memory)</code>	Create GPU buffer, returns handle
<code>gpu.buffer_upload(handle, float_array)</code>	Upload data to buffer
<code>gpu.buffer_download(handle)</code>	Download data from buffer
<code>gpu.create_image(w, h, d, format, usage)</code>	Create GPU image with auto view
<code>gpu.create_sampler(mag, min, address)</code>	Create texture sampler
<code>gpu.load_shader(path, stage)</code>	Load SPIR-V shader module

Descriptors

Function	Description
<code>gpu.create_descriptor_layout(bindings)</code>	Create from array of binding dicts
<code>gpu.create_descriptor_pool(max_sets, sizes)</code>	Create descriptor pool
<code>gpu.allocate_descriptor_set(pool, layout)</code>	Allocate a descriptor set
<code>gpu.update_descriptor(set, binding, type, resource)</code>	Bind buffer/image to descriptor
<code>gpu.update_descriptor_image(set, binding, image, sampler)</code>	Bind combined image sampler

Pipelines

Function	Description
<code>gpu.create_pipeline_layout(layouts, push_size?, stages?)</code>	Create pipeline layout

Function	Description
<code>gpu.create_compute_pipeline(layout, shader)</code>	Create compute pipeline
<code>gpu.create_graphics_pipeline(config_dict)</code>	Create graphics pipeline (see below)
<code>gpu.create_render_pass(attachments)</code>	Create render pass
<code>gpu.create_framebuffer(render_pass, images, w, h)</code>	Create framebuffer

Graphics pipeline config dict keys: - `layout`, `render_pass`, `vertex_shader`, `fragment_shader` (required) - `topology` (default: `TOPO_TRIANGLE_LIST`) - `cull_mode`, `front_face`, `polygon_mode` - `depth_test`, `depth_write` (booleans) - `blend` (boolean, enables alpha blending) - `vertex_bindings`, `vertex_attribs` (arrays of dicts)

Commands

Function	Description
<code>gpu.create_command_pool()</code>	Create command pool
<code>gpu.create_command_buffer(pool)</code>	Allocate command buffer
<code>gpu.begin_commands(cmd)</code>	Begin recording
<code>gpu.end_commands(cmd)</code>	End recording
<code>gpu.cmd_bind_compute_pipeline(cmd, pipe)</code>	Bind compute pipeline
<code>gpu.cmd_bind_graphics_pipeline(cmd, pipe)</code>	Bind graphics pipeline
<code>gpu.cmd_bind_descriptor_set(cmd, layout, index, set)</code>	Bind descriptor set
<code>gpu.cmd_dispatch(cmd, x, y, z)</code>	Dispatch compute workgroups
<code>gpu.cmd_push_constants(cmd, layout, stage, data)</code>	Push constants (float array)
<code>gpu.cmd_begin_render_pass(cmd, rp, fb, clear)</code>	Begin render pass
<code>gpu.cmd_end_render_pass(cmd)</code>	End render pass
<code>gpu.cmd_draw(cmd, verts, instances, first_v, first_i)</code>	Draw call
<code>gpu.cmd_draw_indexed(cmd, indices, instances, first, offset, first_i)</code>	Indexed draw
<code>gpu.cmd_bind_vertex_buffer(cmd, buffer)</code>	Bind vertex buffer
<code>gpu.cmd_set_viewport(cmd, x, y, w, h, min_d, max_d)</code>	Set viewport
<code>gpu.cmd_set_scissor(cmd, x, y, w, h)</code>	Set scissor rect
<code>gpu.cmd_pipeline_barrier(cmd, src, dst, src_acc, dst_acc)</code>	Memory barrier
<code>gpu.cmd_image_barrier(cmd, img, old, new, src, dst, s_acc, d_acc)</code>	Image barrier

Synchronization

Function	Description
<code>gpu.create_fence(signaled?)</code>	Create fence

Function	Description
<code>gpu.wait_fence(fence, timeout?)</code>	Wait for fence
<code>gpu.reset_fence(fence)</code>	Reset fence
<code>gpu.create_semaphore()</code>	Create semaphore
<code>gpu.submit(cmd, wait?, signal?, fence?)</code>	Submit to graphics queue
<code>gpu.submit_compute(cmd, wait?, signal?, fence?)</code>	Submit to compute queue
<code>gpu.submit_with_sync(cmd, wait_sem, signal_sem, fence)</code>	Full sync submit
<code>gpu.device_wait_idle()</code>	Wait for GPU idle

Key Constants

Buffer usage (bitwise OR)

```
gpu.BUFFER_STORAGE  gpu.BUFFER_UNIFORM  gpu.BUFFER_VERTEX
gpu.BUFFER_INDEX    gpu.BUFFER_STAGING  gpu.BUFFER_TRANSFER_SRC
```

Memory

```
gpu.MEMORY_DEVICE_LOCAL  gpu.MEMORY_HOST_VISIBLE  gpu.MEMORY_HOST_COHERENT
```

Formats

```
gpu.FORMAT_RGBA8  gpu.FORMAT_RGBA16F  gpu.FORMAT_RGBA32F
gpu.FORMAT_R32F    gpu.FORMAT_DEPTH32F  gpu.FORMAT_SWAPCHAIN
```

Shader stages

```
gpu.STAGE_VERTEX  gpu.STAGE_FRAGMENT  gpu.STAGE_COMPUTE  gpu.STAGE_ALL
```

Descriptor types

```
gpu.DESC_STORAGE_BUFFER  gpu.DESC_UNIFORM_BUFFER
gpu.DESC_SAMPLED_IMAGE   gpu.DESC_STORAGE_IMAGE  gpu.DESC_COMBINED_SAMPLER
```

Topology

```
gpu.TOPO_TRIANGLE_LIST  gpu.TOPO_LINE_LIST  gpu.TOPO_POINT_LIST
```

Pipeline stages (barriers)

```
gpu.PIPE_TOP  gpu.PIPE_COMPUTE  gpu.PIPE_TRANSFER  gpu.PIPE_FRAGMENT
```

Access flags (barriers)

```
gpu.ACCESS_SHADER_READ  gpu.ACCESS_SHADER_WRITE
gpu.ACCESS_TRANSFER_READ  gpu.ACCESS_HOST_READ
```

Demos

Three demo programs are included in `examples/`:

Demo	File	Description
Empty Window	<code>examples/gpu_window.sage</code>	Creates a window with cycling clear color
Triangle	<code>examples/gpu_triangle.sage</code>	Classic colored triangle via vertex/fragment shaders
3D Hello World	<code>examples/gpu_hello3d.sage</code>	Rotating “HELLO WORLD” text rendered as 3D line segments

Run them with:

```
./sage examples/gpu_window.sage
./sage examples/gpu_triangle.sage
./sage examples/gpu_hello3d.sage
```

Shaders are pre-compiled SPIR-V files in `examples/shaders/`. Recompile with:

```
cd examples/shaders
glslc triangle.vert -o triangle.vert.spv
glslc triangle.frag -o triangle.frag.spv
glslc text3d.vert -o text3d.vert.spv
glslc text3d.frag -o text3d.frag.spv
```

9.21 Hardware Natives (Embedded Targets)

For embedded targets (such as the Pico or other bare-metal environments), the SageLang C backend provides native hardware interaction routines through the `hw` namespace. These bindings map directly to standard C hardware APIs.

Function Signature	Description
<code>hw.gpio_init(pin)</code>	Initializes a GPIO pin.
<code>hw.gpio_set_dir(pin, out)</code>	Sets the direction of a GPIO pin (e.g., input or output).
<code>hw.gpio_put(pin, val)</code>	Sets the output value of a GPIO pin.
<code>hw.gpio_get(pin)</code>	Reads the value of a GPIO pin.
<code>hw.gpio_set_pull(pin, up, down)</code>	Configures pull-up/pull-down resistors for a pin.
<code>hw.clock_hz()</code>	Returns the system clock frequency in Hz.
<code>hw.uptime_ms()</code>	Returns the system uptime in milliseconds.
<code>hw.delay_ms(ms)</code>	Blocks execution for <code>ms</code> milliseconds.
<code>hw.delay_us(us)</code>	Blocks execution for <code>us</code> microseconds.
<code>hw.uart_init(baud)</code>	Initializes the default UART interface at the specified baud rate.
<code>hw.uart_putc(c)</code>	Transmits a single character over UART.
<code>hw.uart_puts(s)</code>	Transmits a string over UART.
<code>hw.uart_getc()</code>	Reads a single character from UART (returns -1 if no data is available).
<code>hw.adc_init(pin)</code>	Initializes an ADC (Analog-to-Digital Converter) pin.
<code>hw.adc_read(pin)</code>	Reads the analog value from the specified ADC pin.
<code>hw.temp_c()</code>	Returns the internal temperature sensor reading in degrees Celsius.
<code>hw.rgb_set(r, g, b)</code>	Sets the color of an onboard RGB LED (e.g., WS2812).
<code>hw.spi_init(bus, baud)</code>	Initializes an SPI bus.
<code>hw.spi_write(bus, data)</code>	Writes data to the SPI bus.
<code>hw.spi_read(bus, count)</code>	Reads <code>count</code> bytes from the SPI bus.
<code>hw.lcd_fb_init(w, h)</code>	Initializes an LCD framebuffer of width <code>w</code> and height <code>h</code> .

Function Signature	Description
<code>hw.lcd_fb_pixel(x, y, color)</code>	Sets the color of a specific pixel in the LCD framebuffer.
<code>hw.lcd_fb_fill(color)</code>	Fills the entire LCD framebuffer with <code>color</code> .
<code>hw.lcd_fb_flush_bytes(count)</code>	Flushes <code>count</code> bytes of the framebuffer to the LCD screen.

Note: These functions are resolved during compilation by the C backend when targeting appropriate platforms.

Appendix: Quick Reference

Keywords

and as async await break case catch class comptime continue default defer elif else end enum false final

Note: `print`, `end`, `match`, `init`, `enum`, `struct`, and `trait` are soft keywords, meaning they can be used as variable, property, or method names where unambiguous.

Built-in Functions

```
len(x) push(arr, val) pop(arr) append(arr, val) range(a, b)
split(str, delim) join(arr, sep) replace(s, old, new)
upper(s) lower(s) strip(s) slice(arr/str, start, end)
str(x) tonumber(s) int(x) input() clock()
type(x) val_tag(x) build_info() chr(n) ord(c) hash(x) sizeof(x) doc(x)
startswith(s, prefix) endswith(s, suffix)
contains(s, sub) indexof(s, sub) string_count(s, sub) string_repeat(s, n)
array_extend(a, b) array_repeat(a, n) array_reverse(a)
array_contains(a, v) array_index_of(a, v) array_min(a) array_max(a)
array_sum(a) array_product(a)
bytes() bytes_len(b) bytes_get(b, i) bytes_set(b, i, v)
bytes_push(b, v) bytes_slice(b, s, e) bytes_to_string(b)
dict_keys(d) dict_values(d) dict_has(d, k) dict_delete(d, k) hasattr(obj, name)
gc_collect() gc_stats() gc_collections() gc_enable() gc_disable()
gc_mode() gc_set_arc() gc_set_orc()
next(gen)
ffi_open(path) ffi_call(lib, fn, ret, ...) ffi_sym(lib, name) ffi_close(lib)
ffi_sym_addr(lib, name)
mem_alloc(n) mem_free(ptr) mem_read(ptr, off, type) mem_write(ptr, off, type, val)
mem_size(ptr) addressof(val) addressof_raw(val) ptr_add(ptr, n) ptr_to_int(ptr)
path_join(a, b) path_dirname(p) path_basename(p) path_ext(p)
path_exists(p) path_is_dir(p) path_is_file(p)
asm_exec(code, ret, ...) asm_compile(code, arch, out) asm_arch()
struct_def(fields) struct_new(def) struct_get(ptr, def, name)
struct_set(ptr, def, name, val) struct_size(def)
cpu_count() cpu_physical_cores() cpu_has_hyperthreading()
thread_set_affinity(id) thread_get_core()
atomic_new(v) atomic_load(a) atomic_store(a, v)
atomic_add(a, v) atomic_cas(a, e, d) atomic_exchange(a, v)
sem_new(n) sem_wait(s) sem_post(s) sem_trywait(s)
vm_gas_limit_set(n) vm_gas_used_get() vm_gas_limit_get()
build_quad_verts(quads) build_line_quads(lines, thickness, r, g, b, a)
```

Operators

Arithmetic:	+ - * / %
Comparison:	== != < > <= >=
Logical:	and or not
Bitwise:	& ^ ~ << >>
Assignment:	=
Indexing:	arr[i]
Slicing:	arr[a:b]
Property:	obj.field or obj->field
Call:	func(args)